

leading practices

leading practices

events should be processed in realtime
(i.e. within a few seconds)

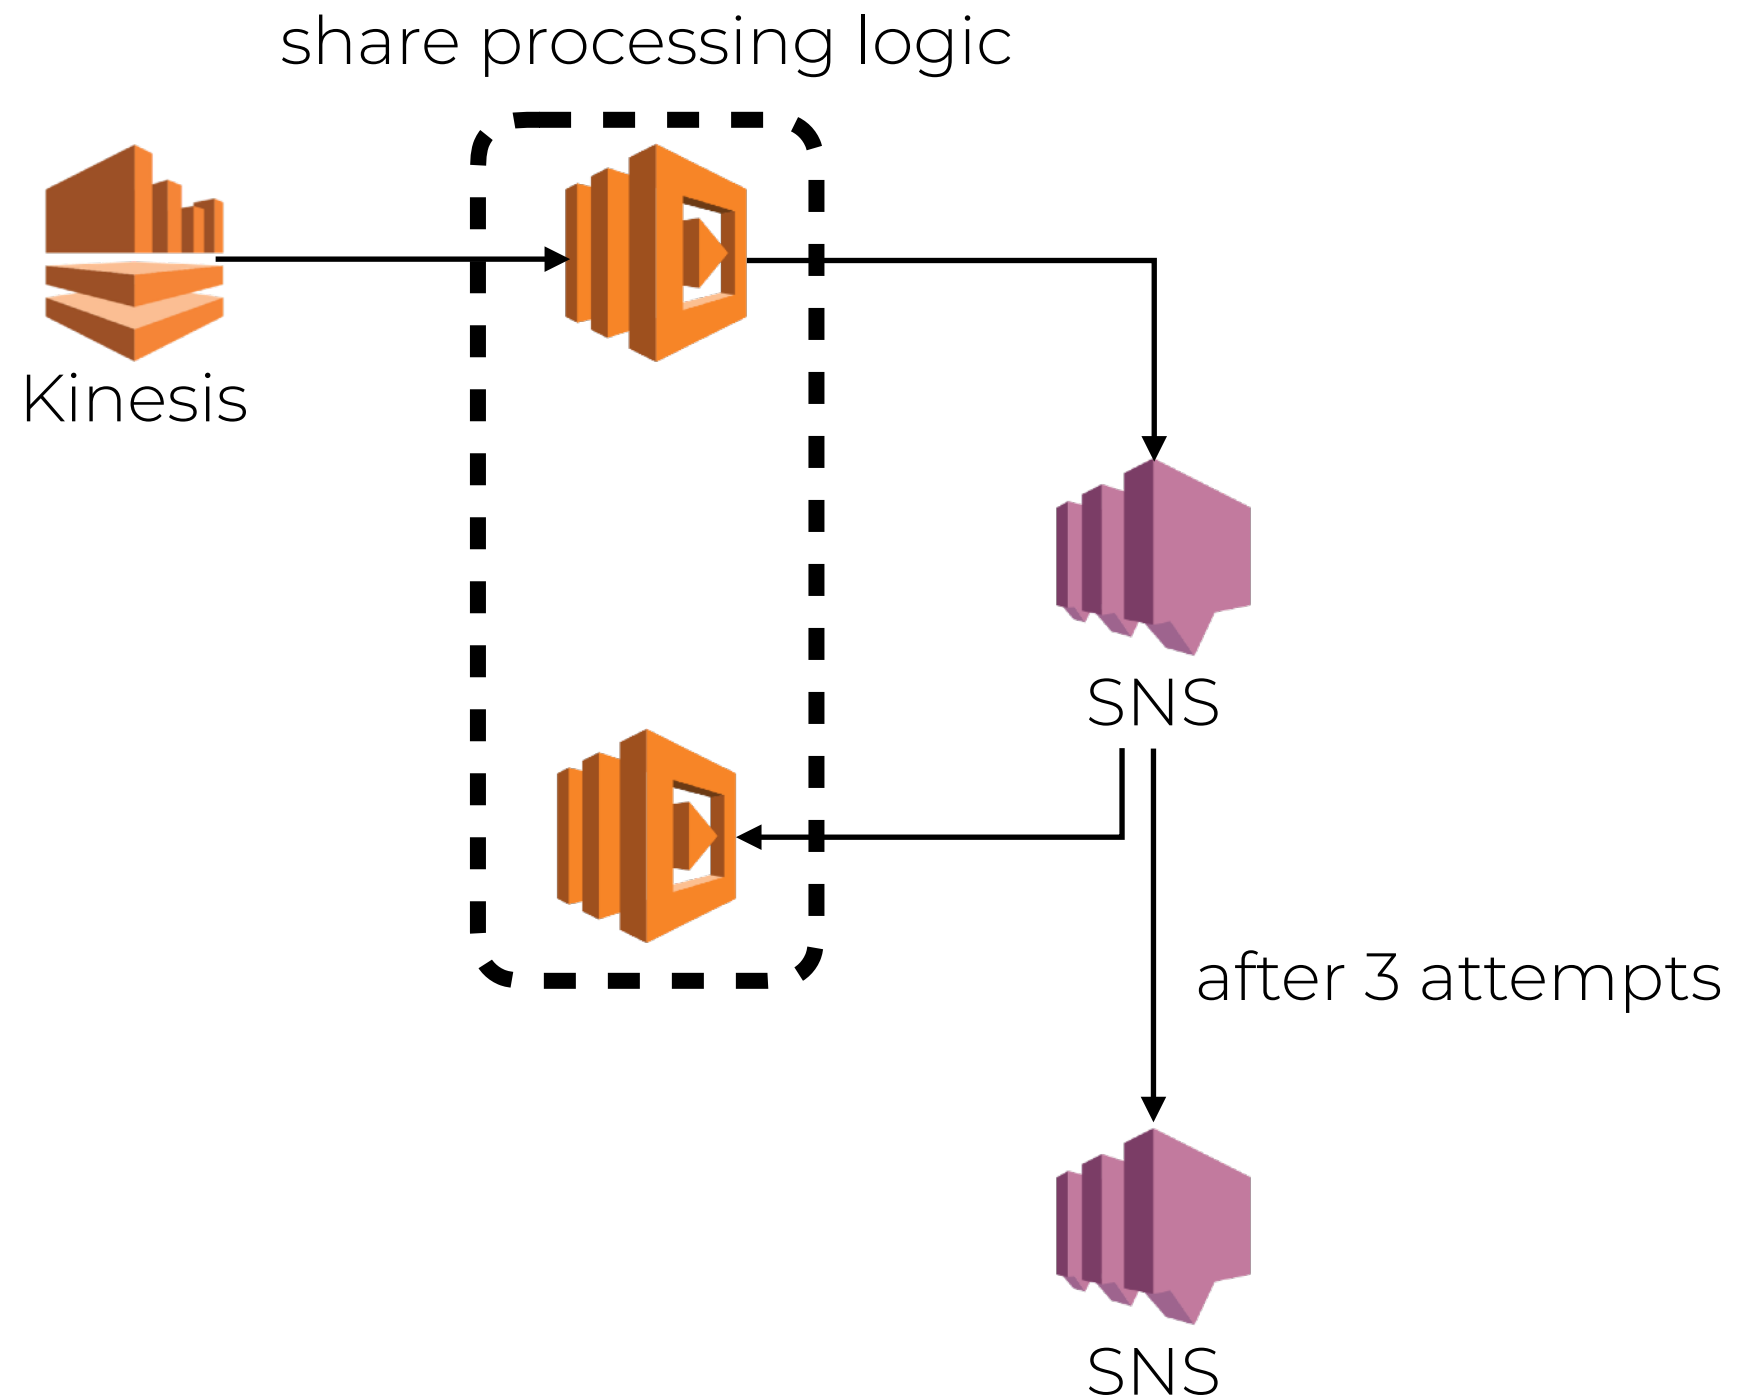
leading practices

failed events should be retried, but the retries
should not violate the realtime constraint

leading practices

unprocessed events should be retrievable
(i.e. available for human intervention, or root cause analysis)

leading practices



leading practices

Put Requests	\$0.014 per mill (at 25KB chunks)
Shard Hours	\$0.015 per hour per shard 1MB/s ingress, 1000 records/s 2MB/s egress, 5 read req/s
Extended Retention	\$0.020 per shard hour

<http://amzn.to/2ubyaot>

leading practices

optimise payload size

Put Requests

\$0.014 per mill (at 25KB chunks)

Shard Hours

\$0.015 per hour per shard
1MB/s ingress, 1000 records/s
2MB/s egress, 5 read req/s

Extended Retention

\$0.020 per shard hour

leading practices

consider using a binary format such as
Protocol Buffer and Thrift

leading practices

alarm on **IteratorAge** and
WriteProvisionedThroughputExceeded



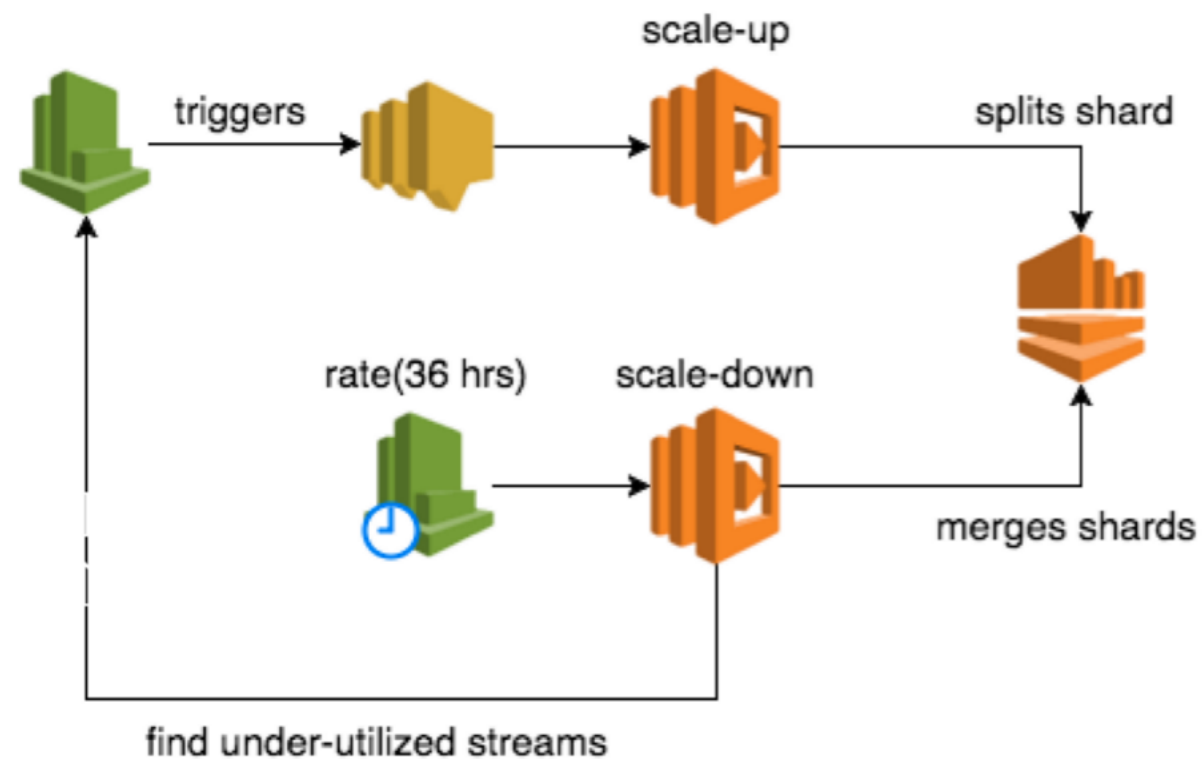
Yan Cui

Senior Developer @ Space Ape Games.

May 5 · 6 min read

Auto-scaling Kinesis streams with AWS Lambda

A recipe for creating a cost-effective solution for auto-scaling Kinesis streams using Lambda functions




<http://bit.ly/2hxZGui>

leading practices

avoid “hot” streams

leading practices





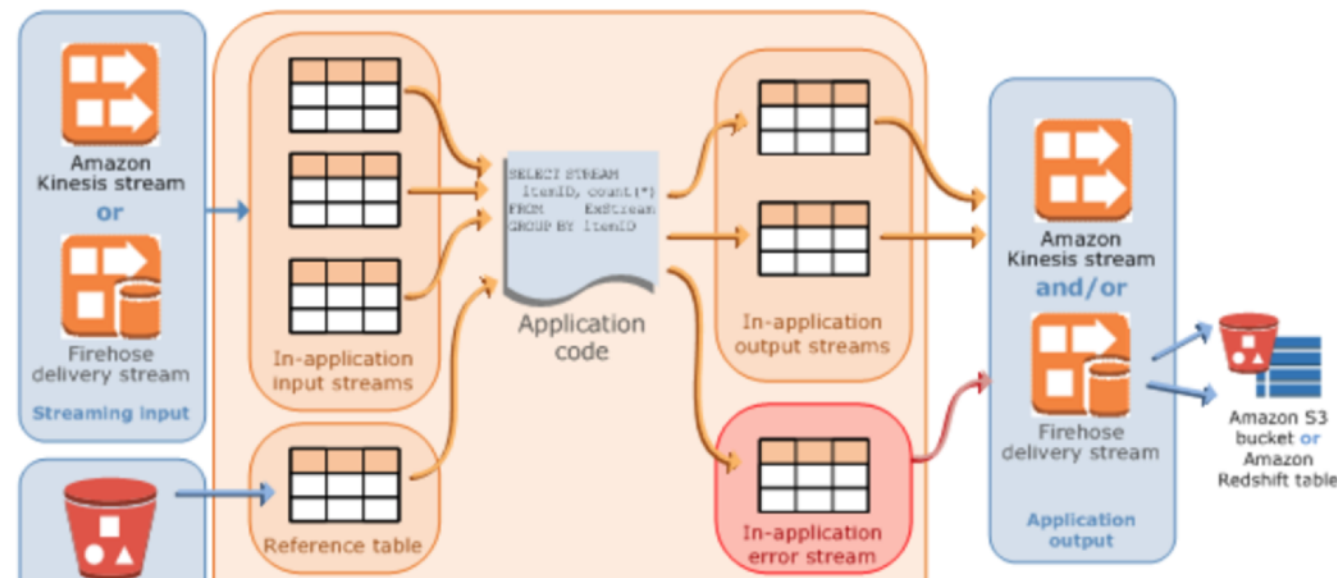
AWS Lambda
docs

If your stream has 100 active shards, there will be at most 100 Lambda function invocations running concurrently. Then, each Lambda function processes events on a shard **in the order that they arrive**.

<http://amzn.to/2va5GN5>

leading practices

Amazon Kinesis Streams fan-out via Kinesis Analytics - made with **serverless** ⚡



Source: <https://github.com/alexcasalboni/kinesis-streams-fan-out-kinesis-analytics>

How to fan-out Amazon Kinesis Streams?

Published on June 23, 2017



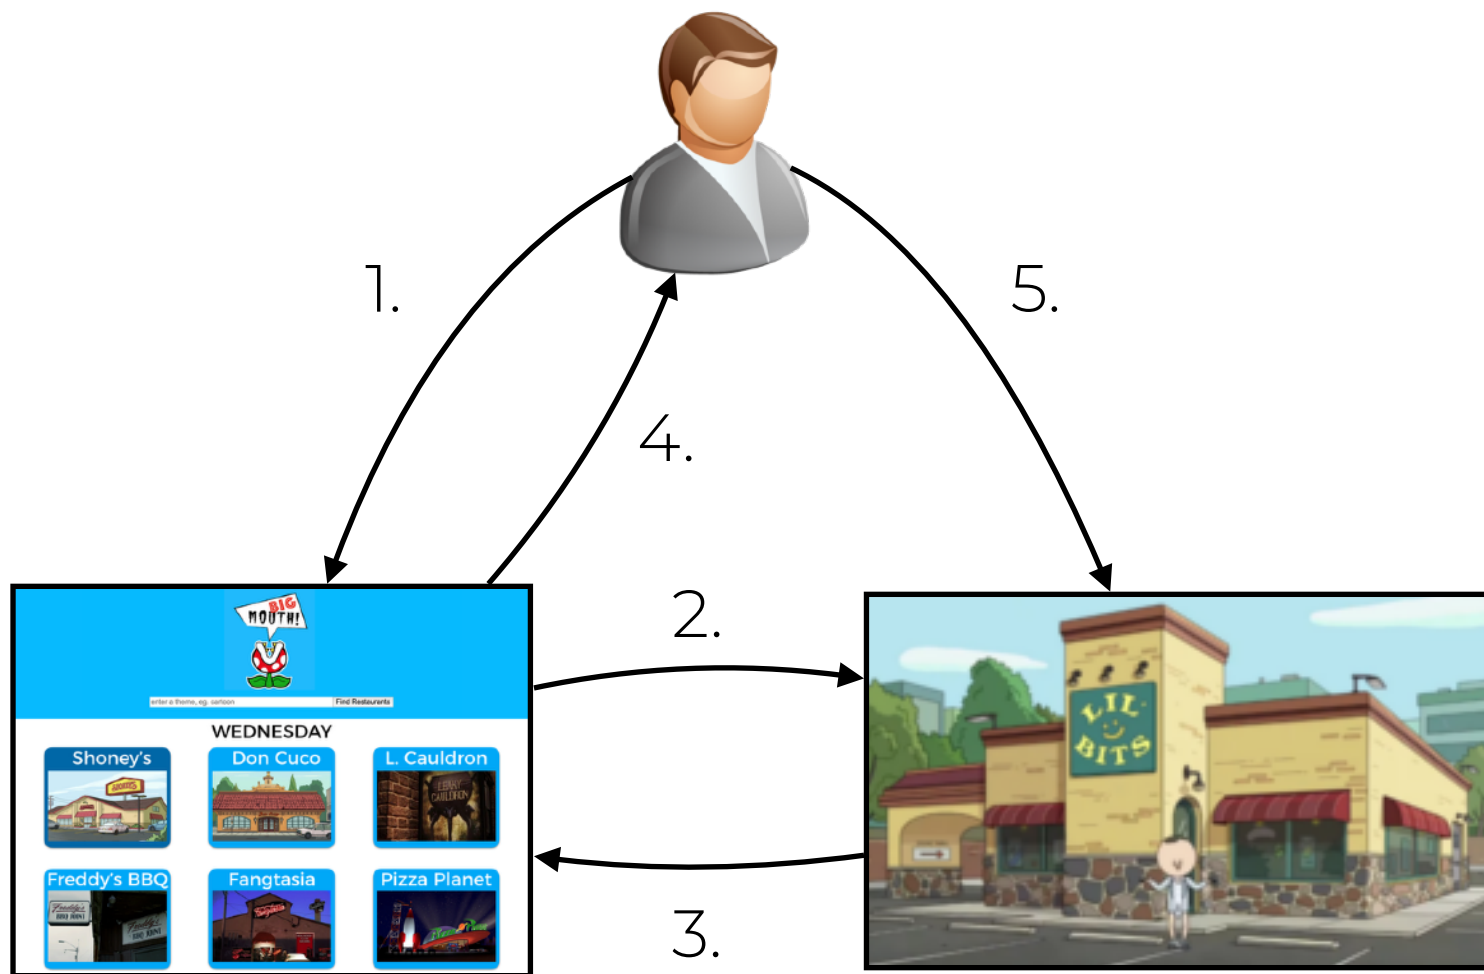
Alex Casalboni | ✓ Following
Cloud Evangelist, Serverless lover, AI enthusiast
[4 articles](#)

👍 26 💬 0 ➦ 23

<http://bit.ly/kinesis-fan-out>


A case against events

order flow



1. customer places order
2. notifies restaurant
3. restaurant accepts order
4. customer confirmation
5. pick up

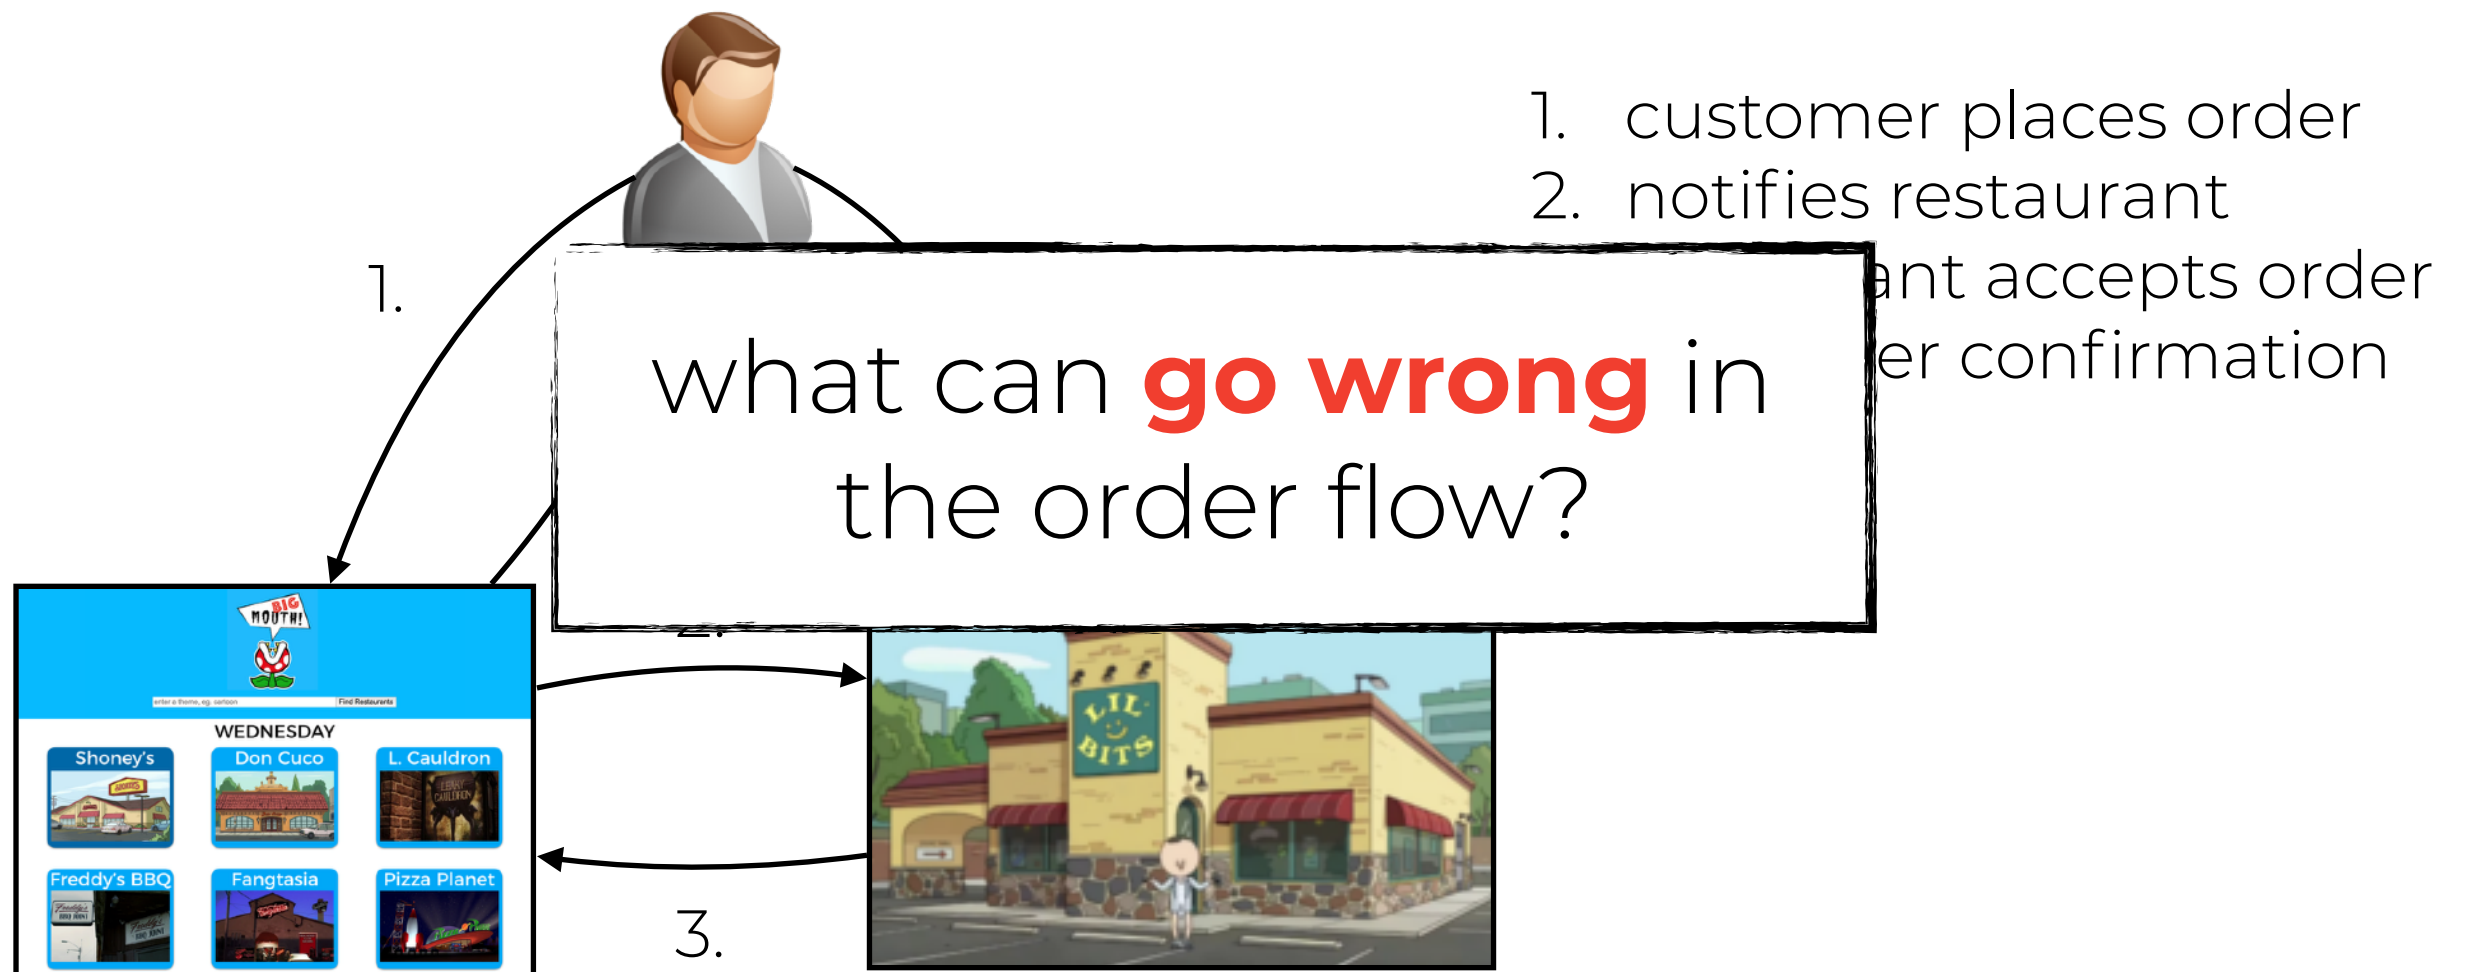
order flow



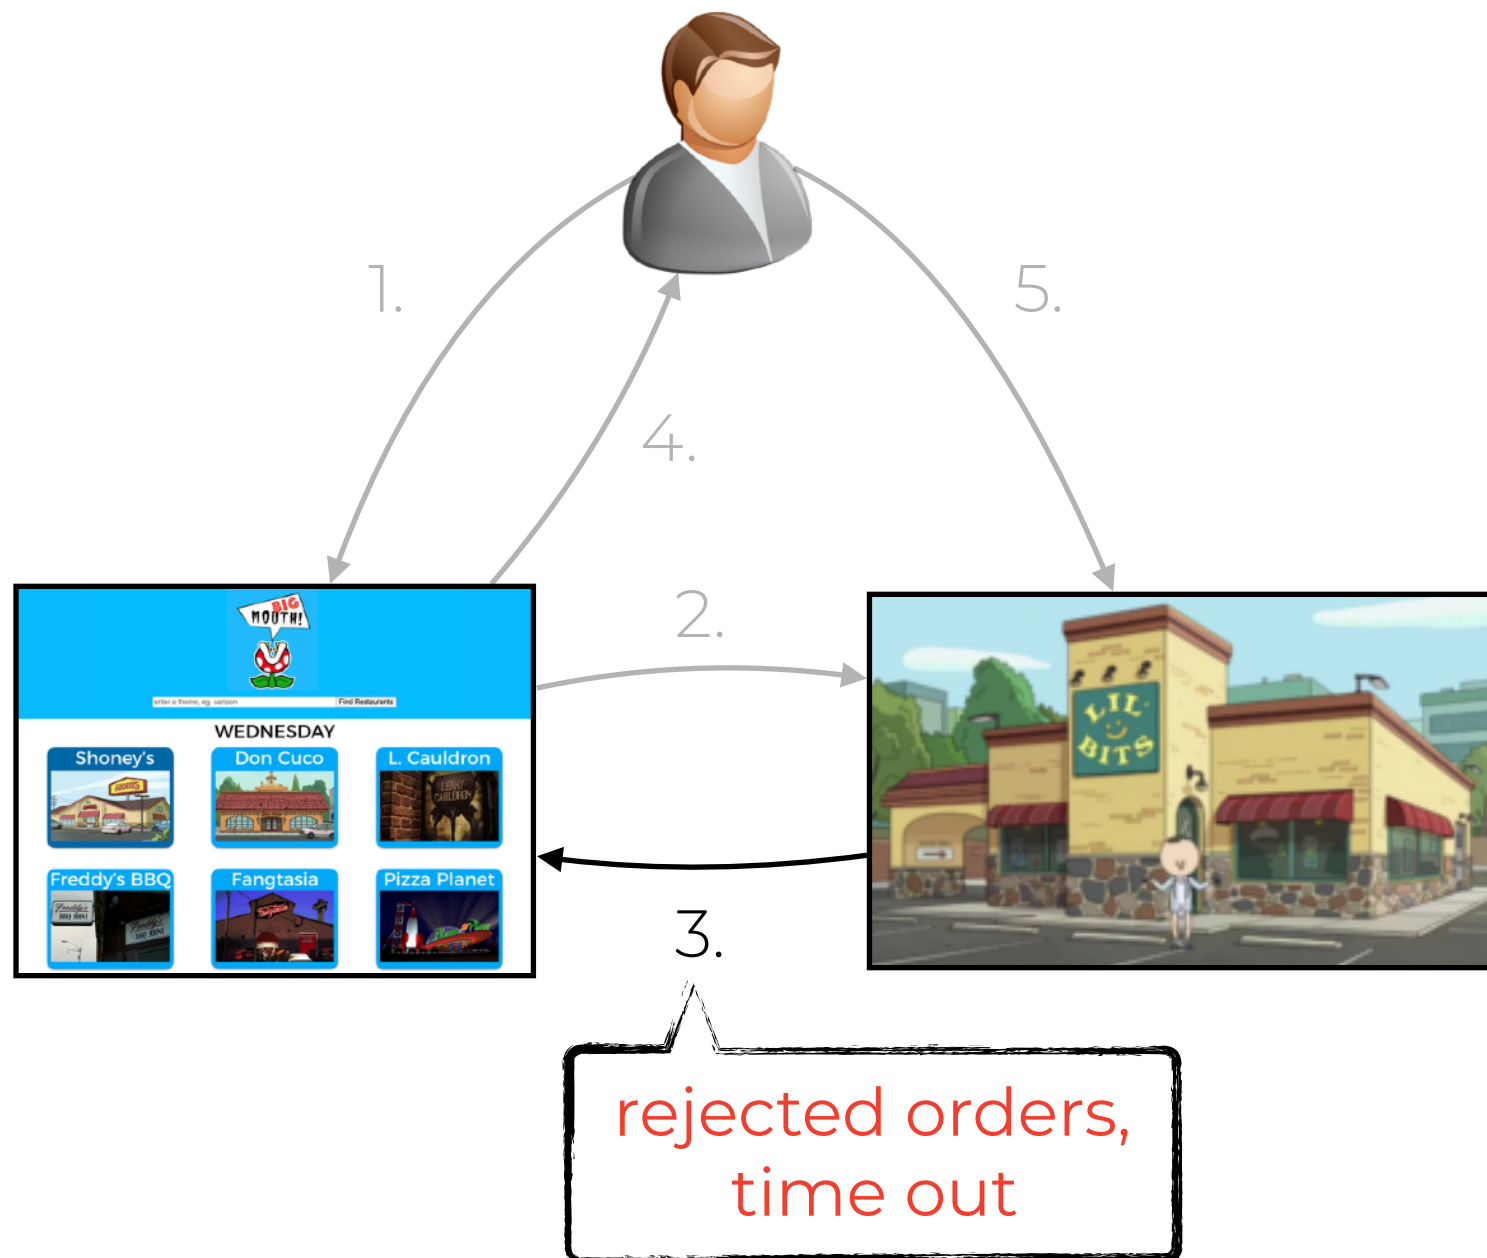
event-driven

- loose coupling
- scalability

failure modes




failure modes



1. customer places order
2. notifies restaurant
3. restaurant accepts order
4. customer confirmation
5. pick up

failure modes

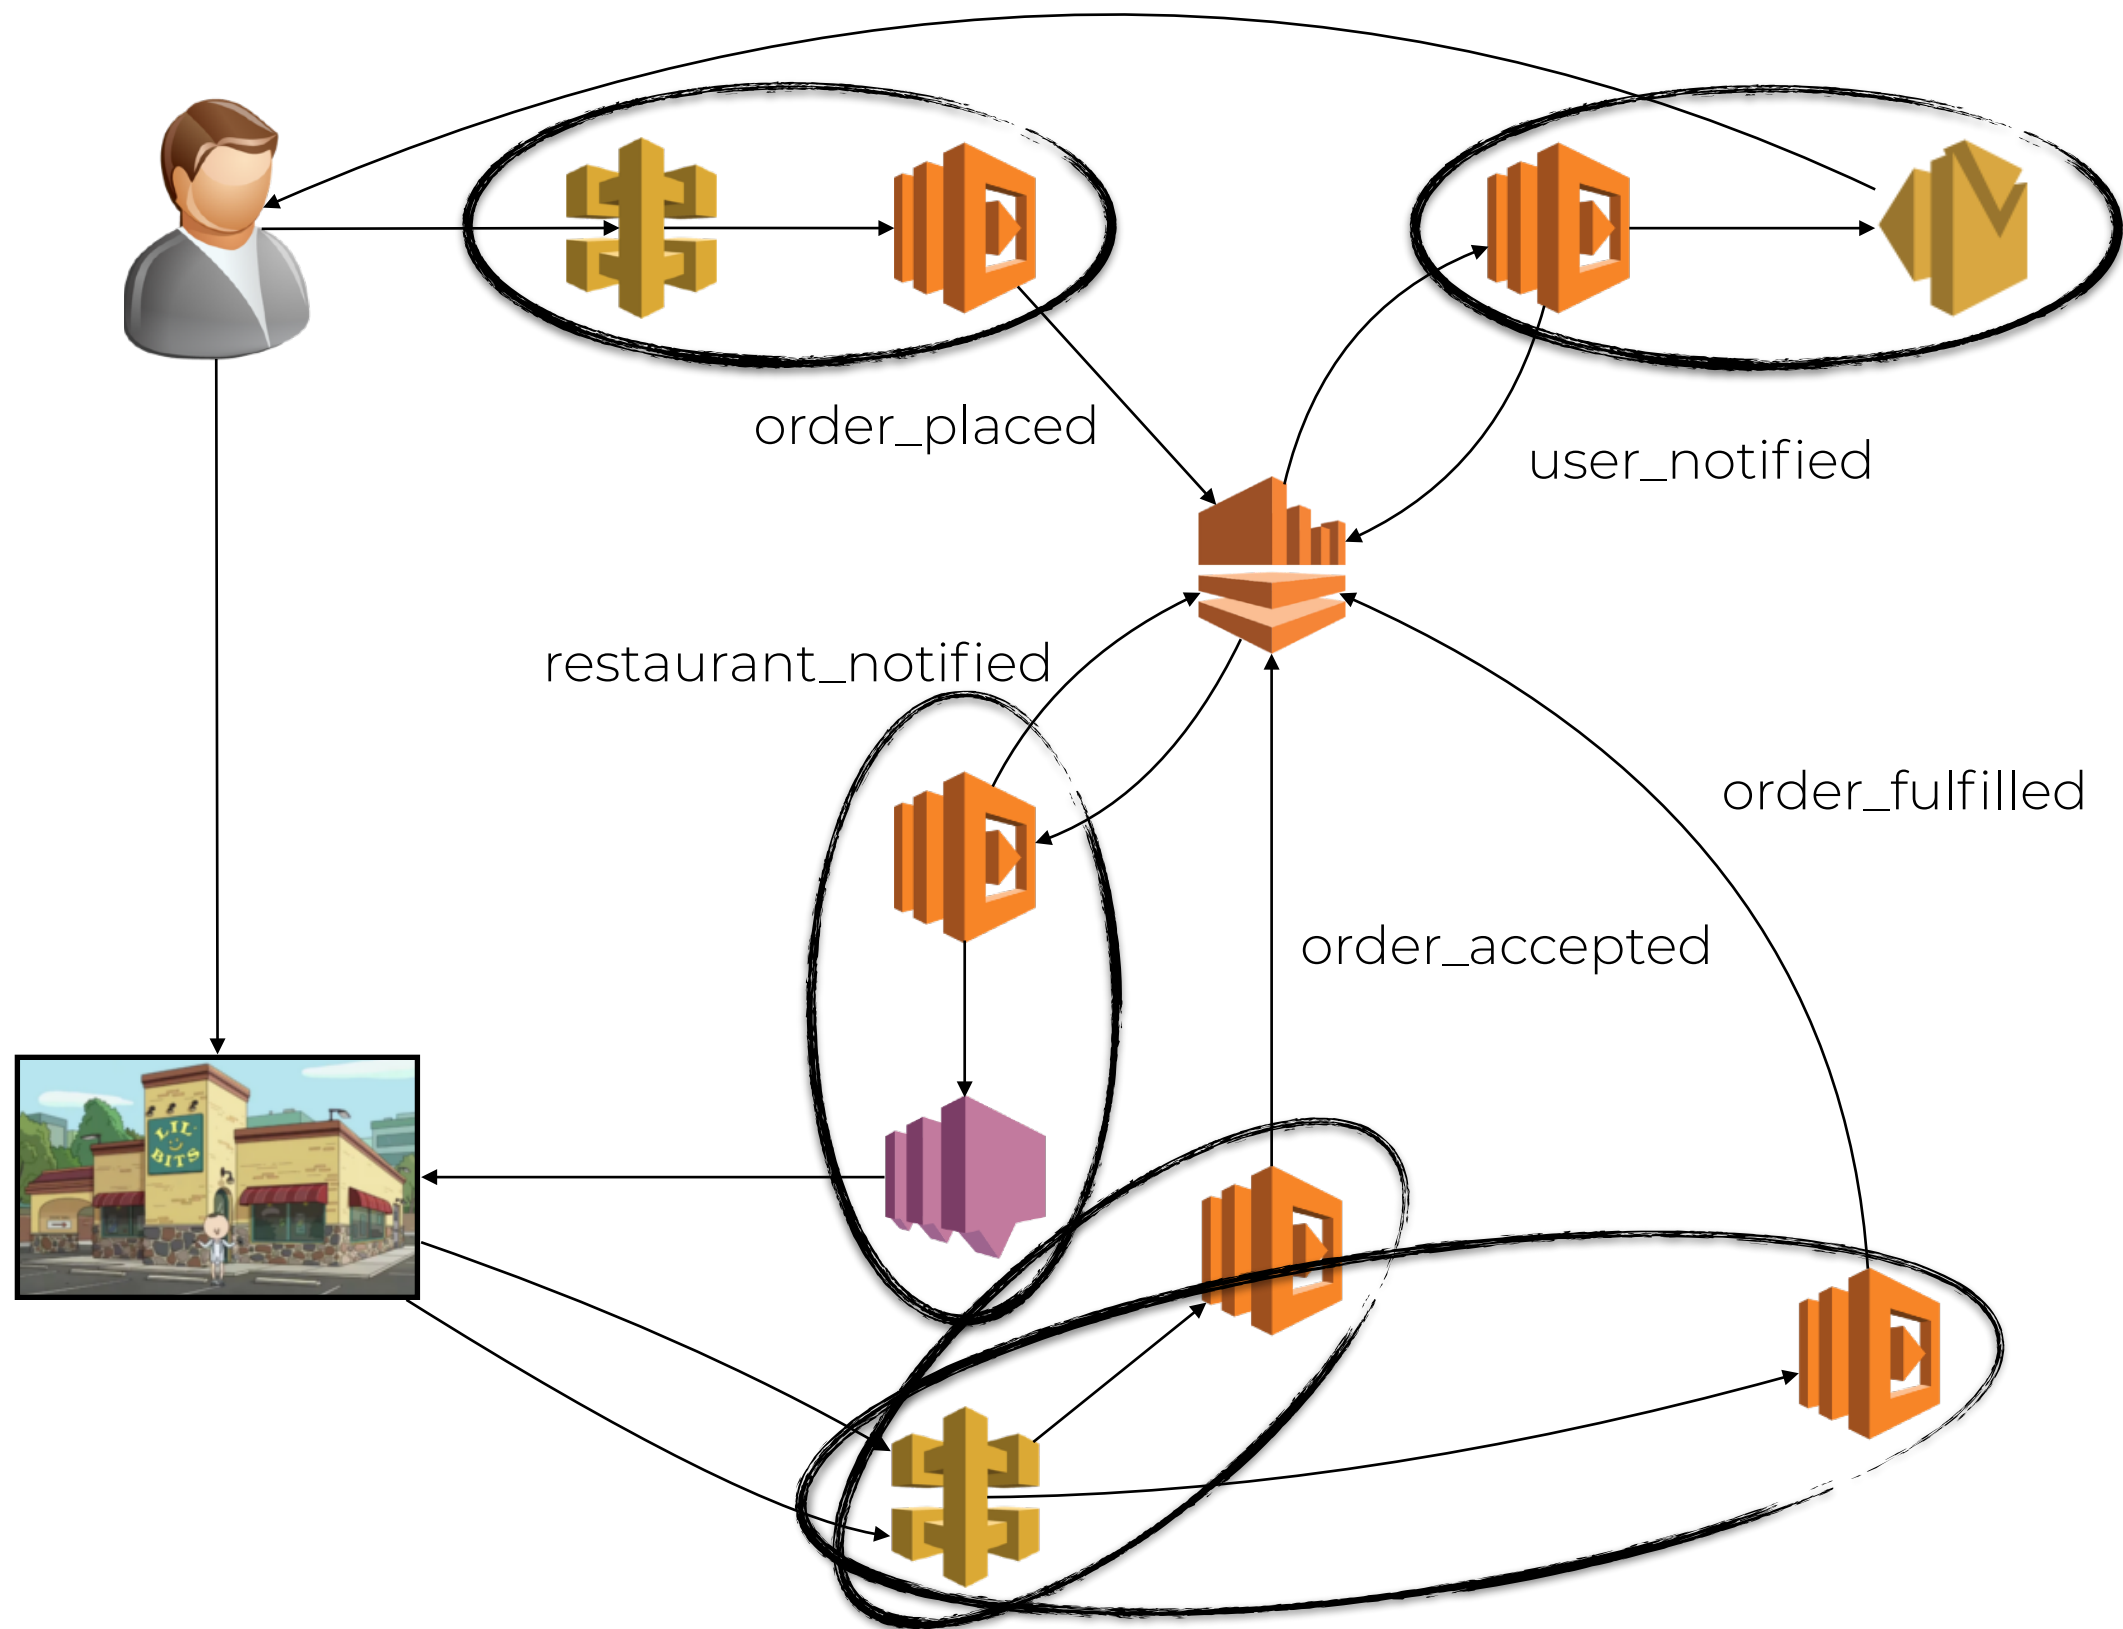


1. customer places order
2. notifies restaurant
3. restaurant accepts order
4. customer confirmation
5. pick up

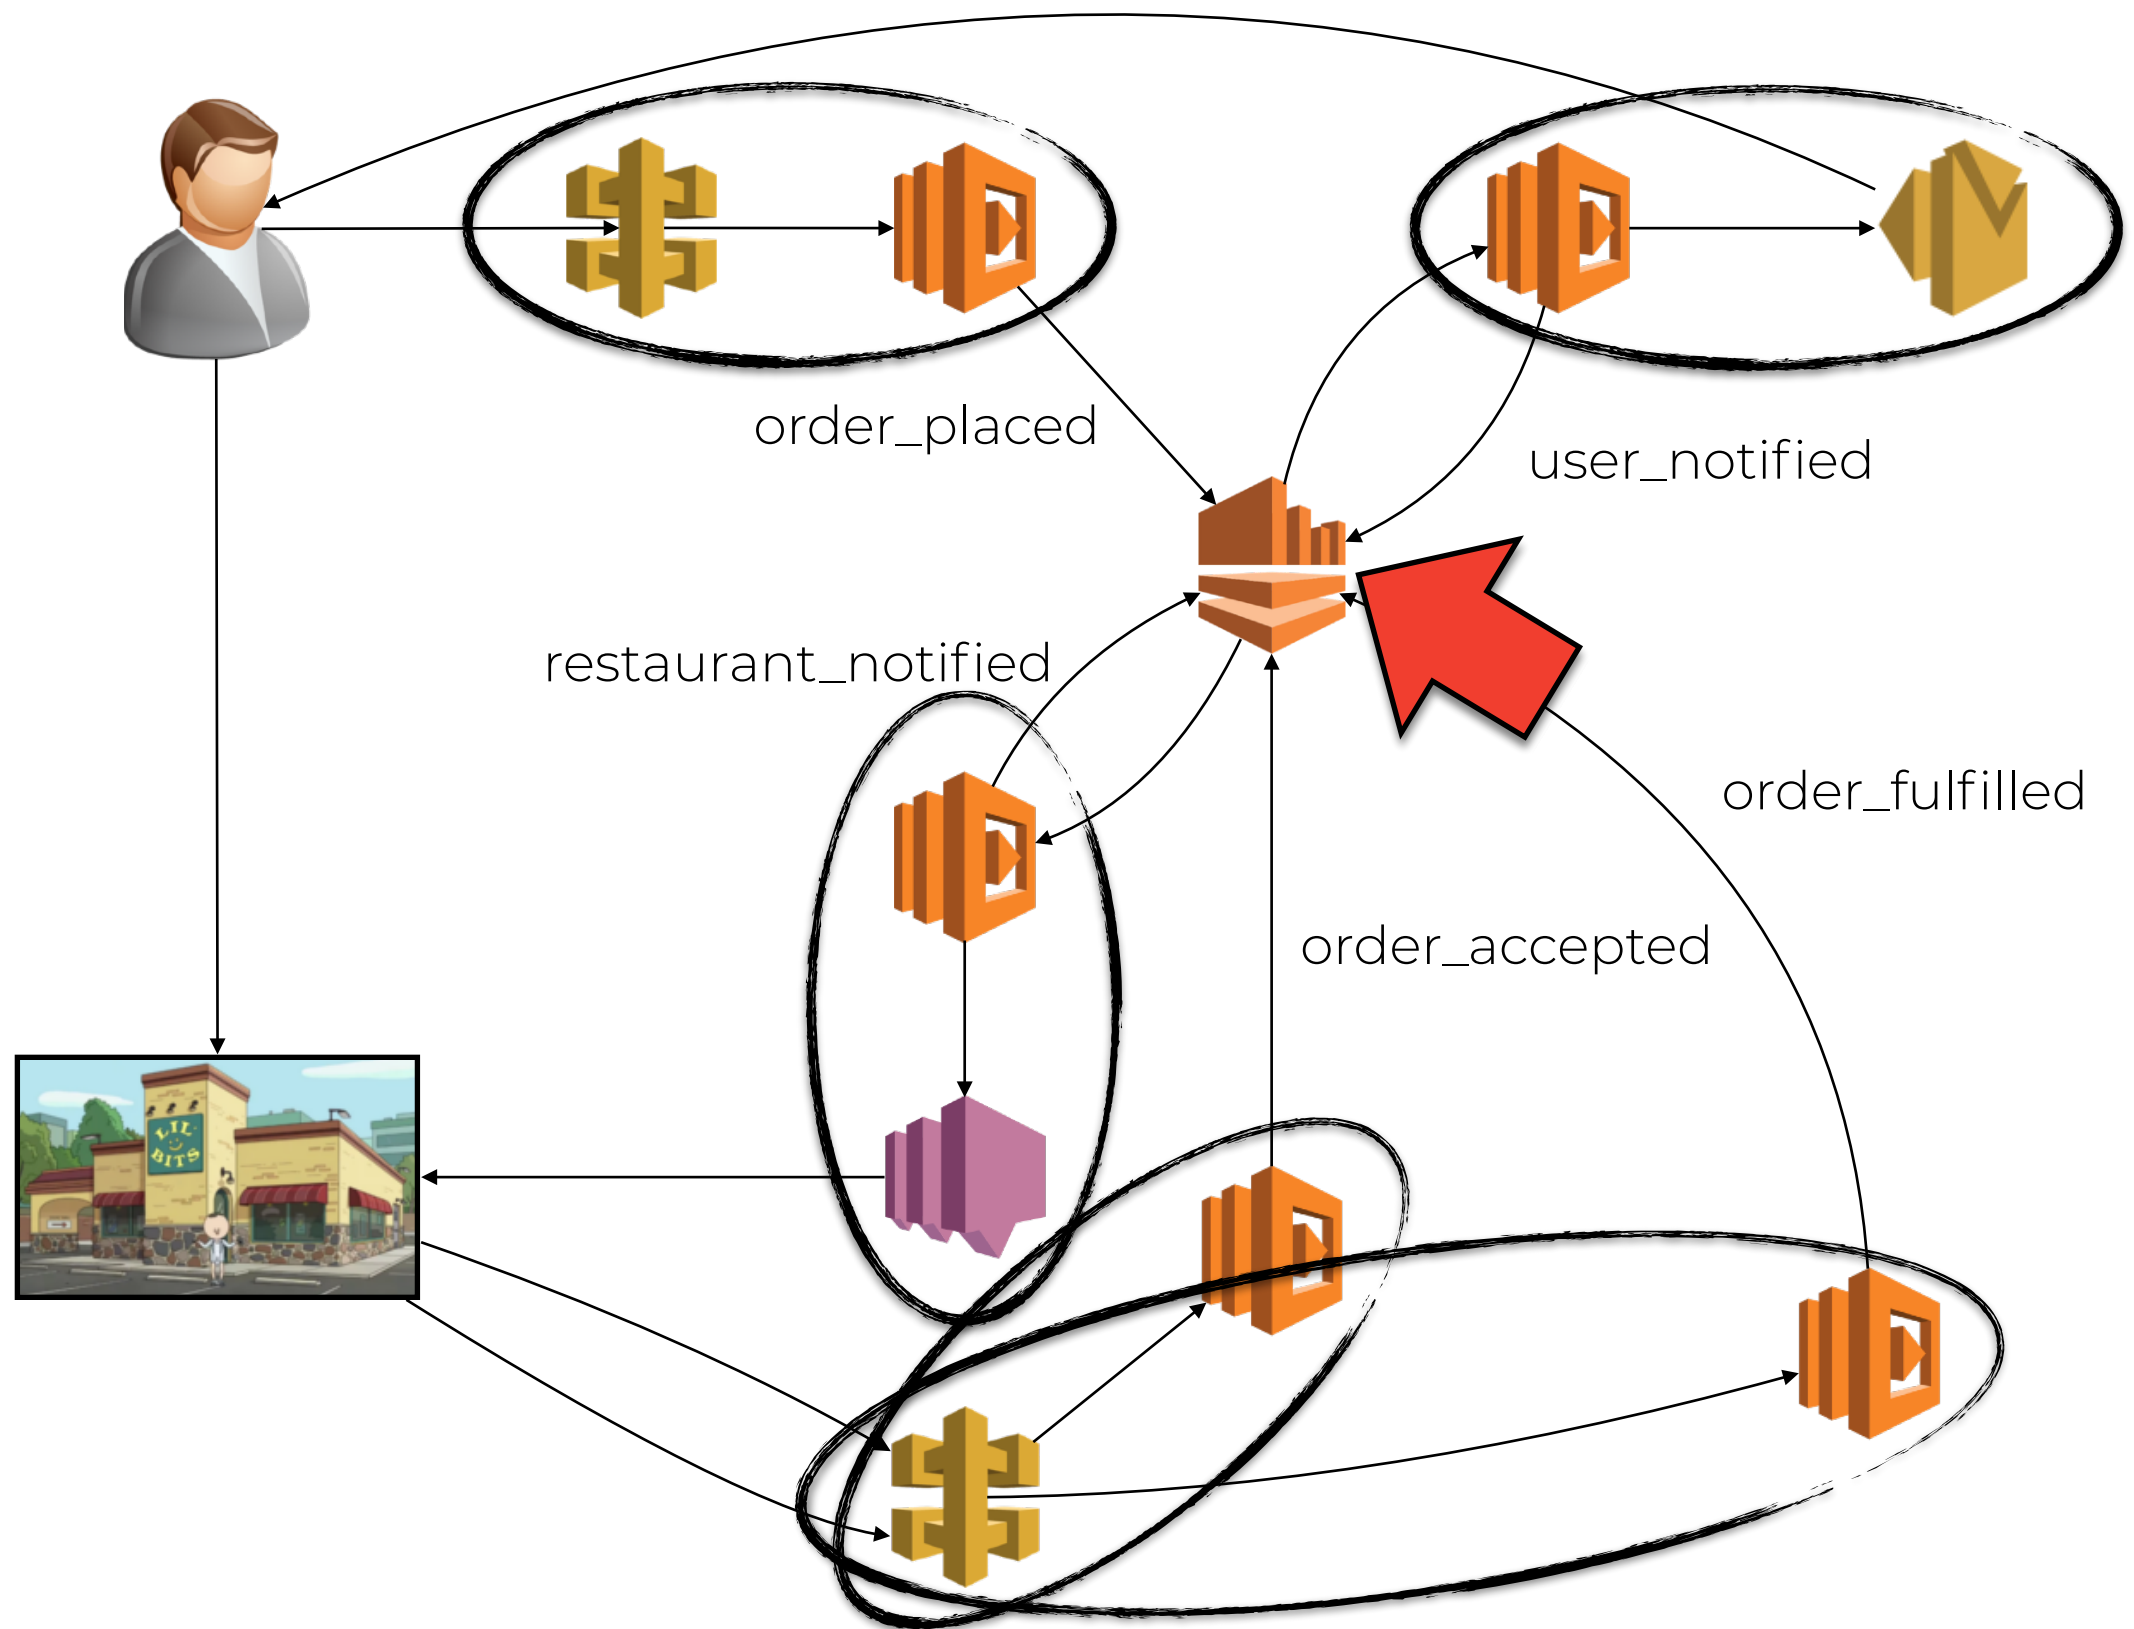
failure modes

- minimise impact on user experience

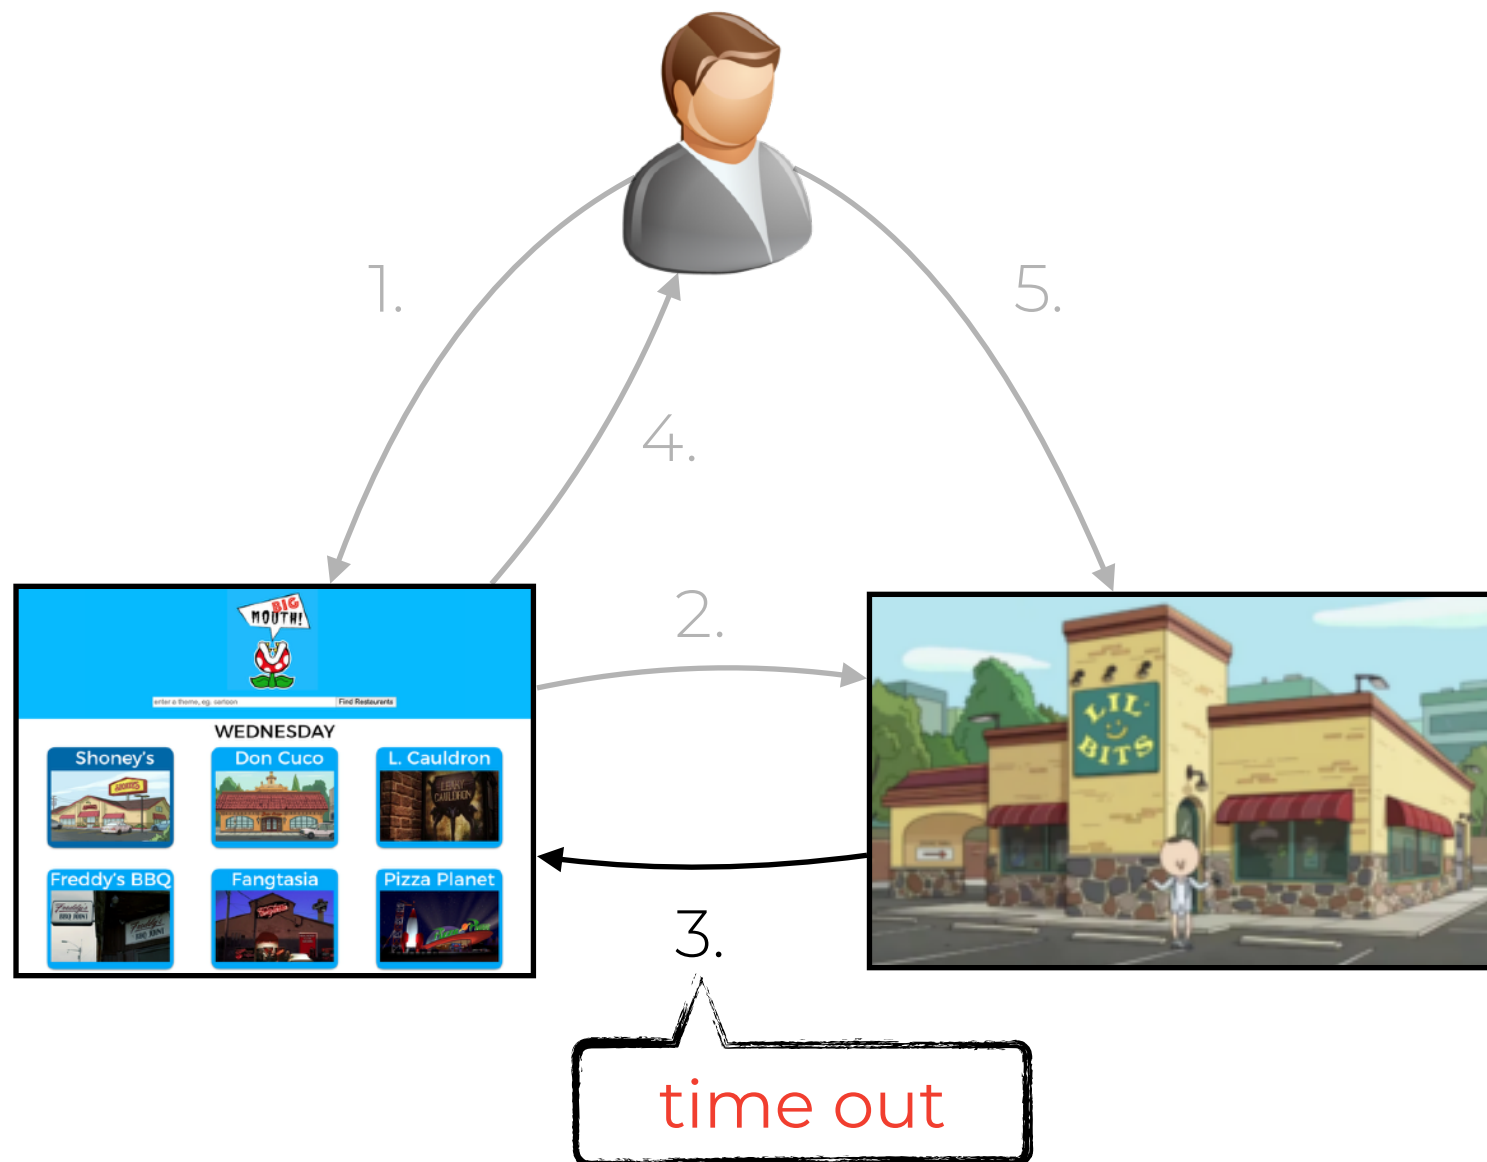
order flow



order flow

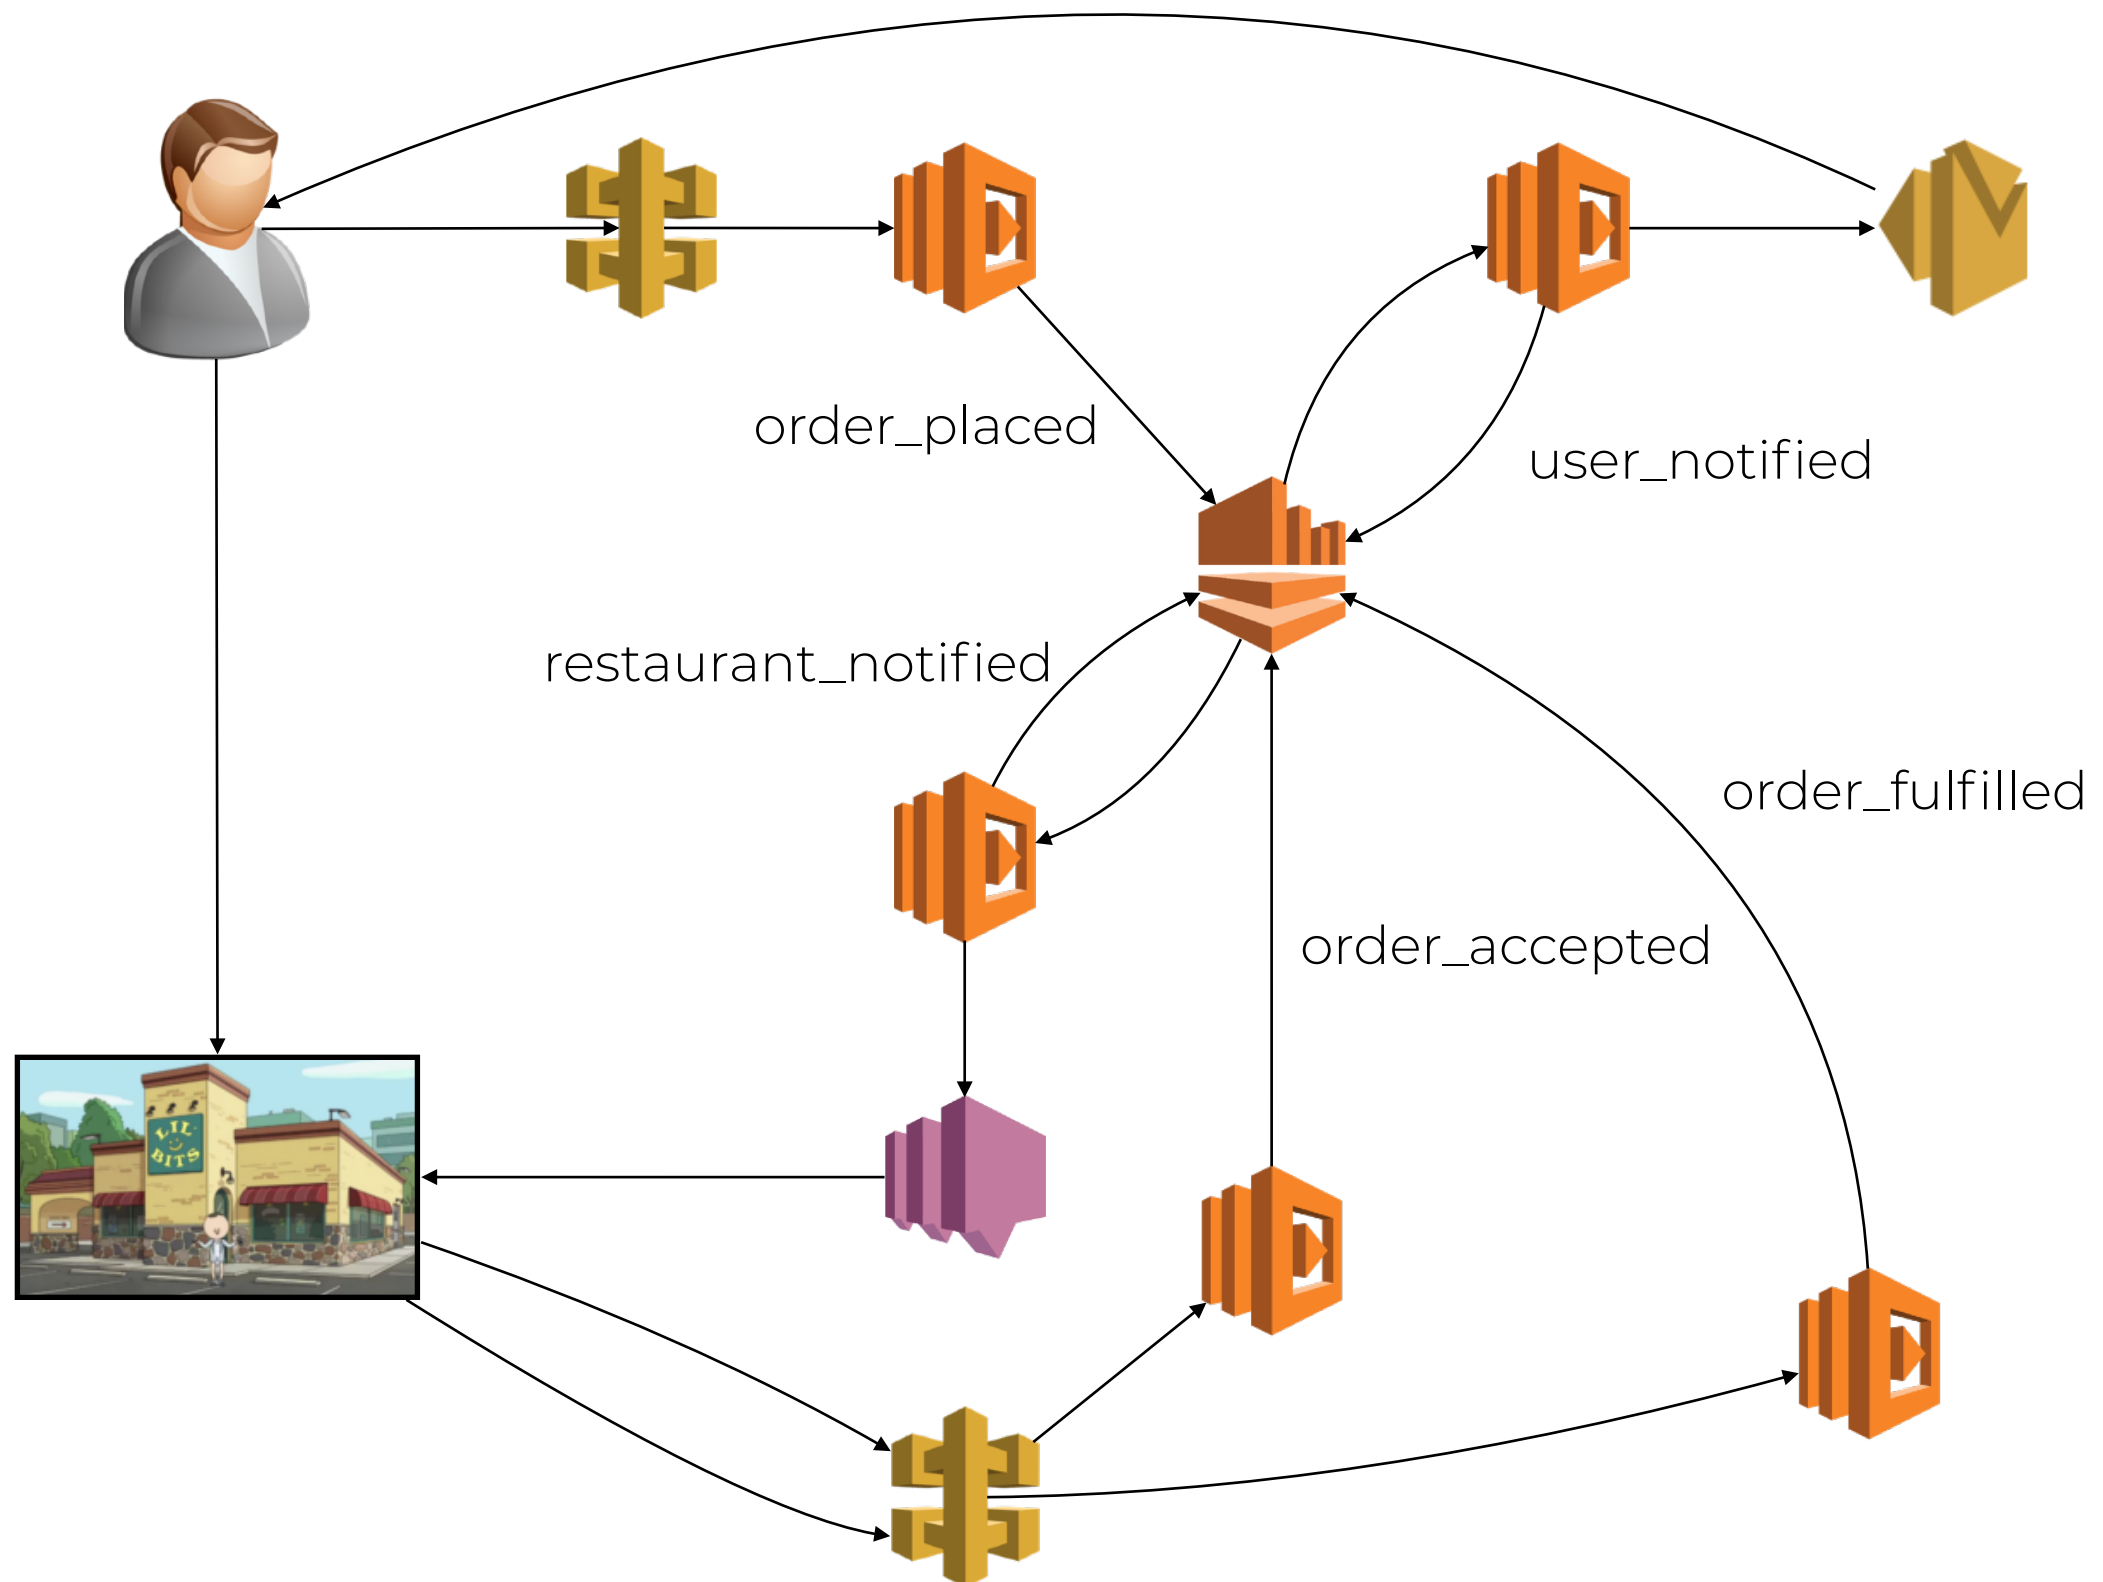


failure modes

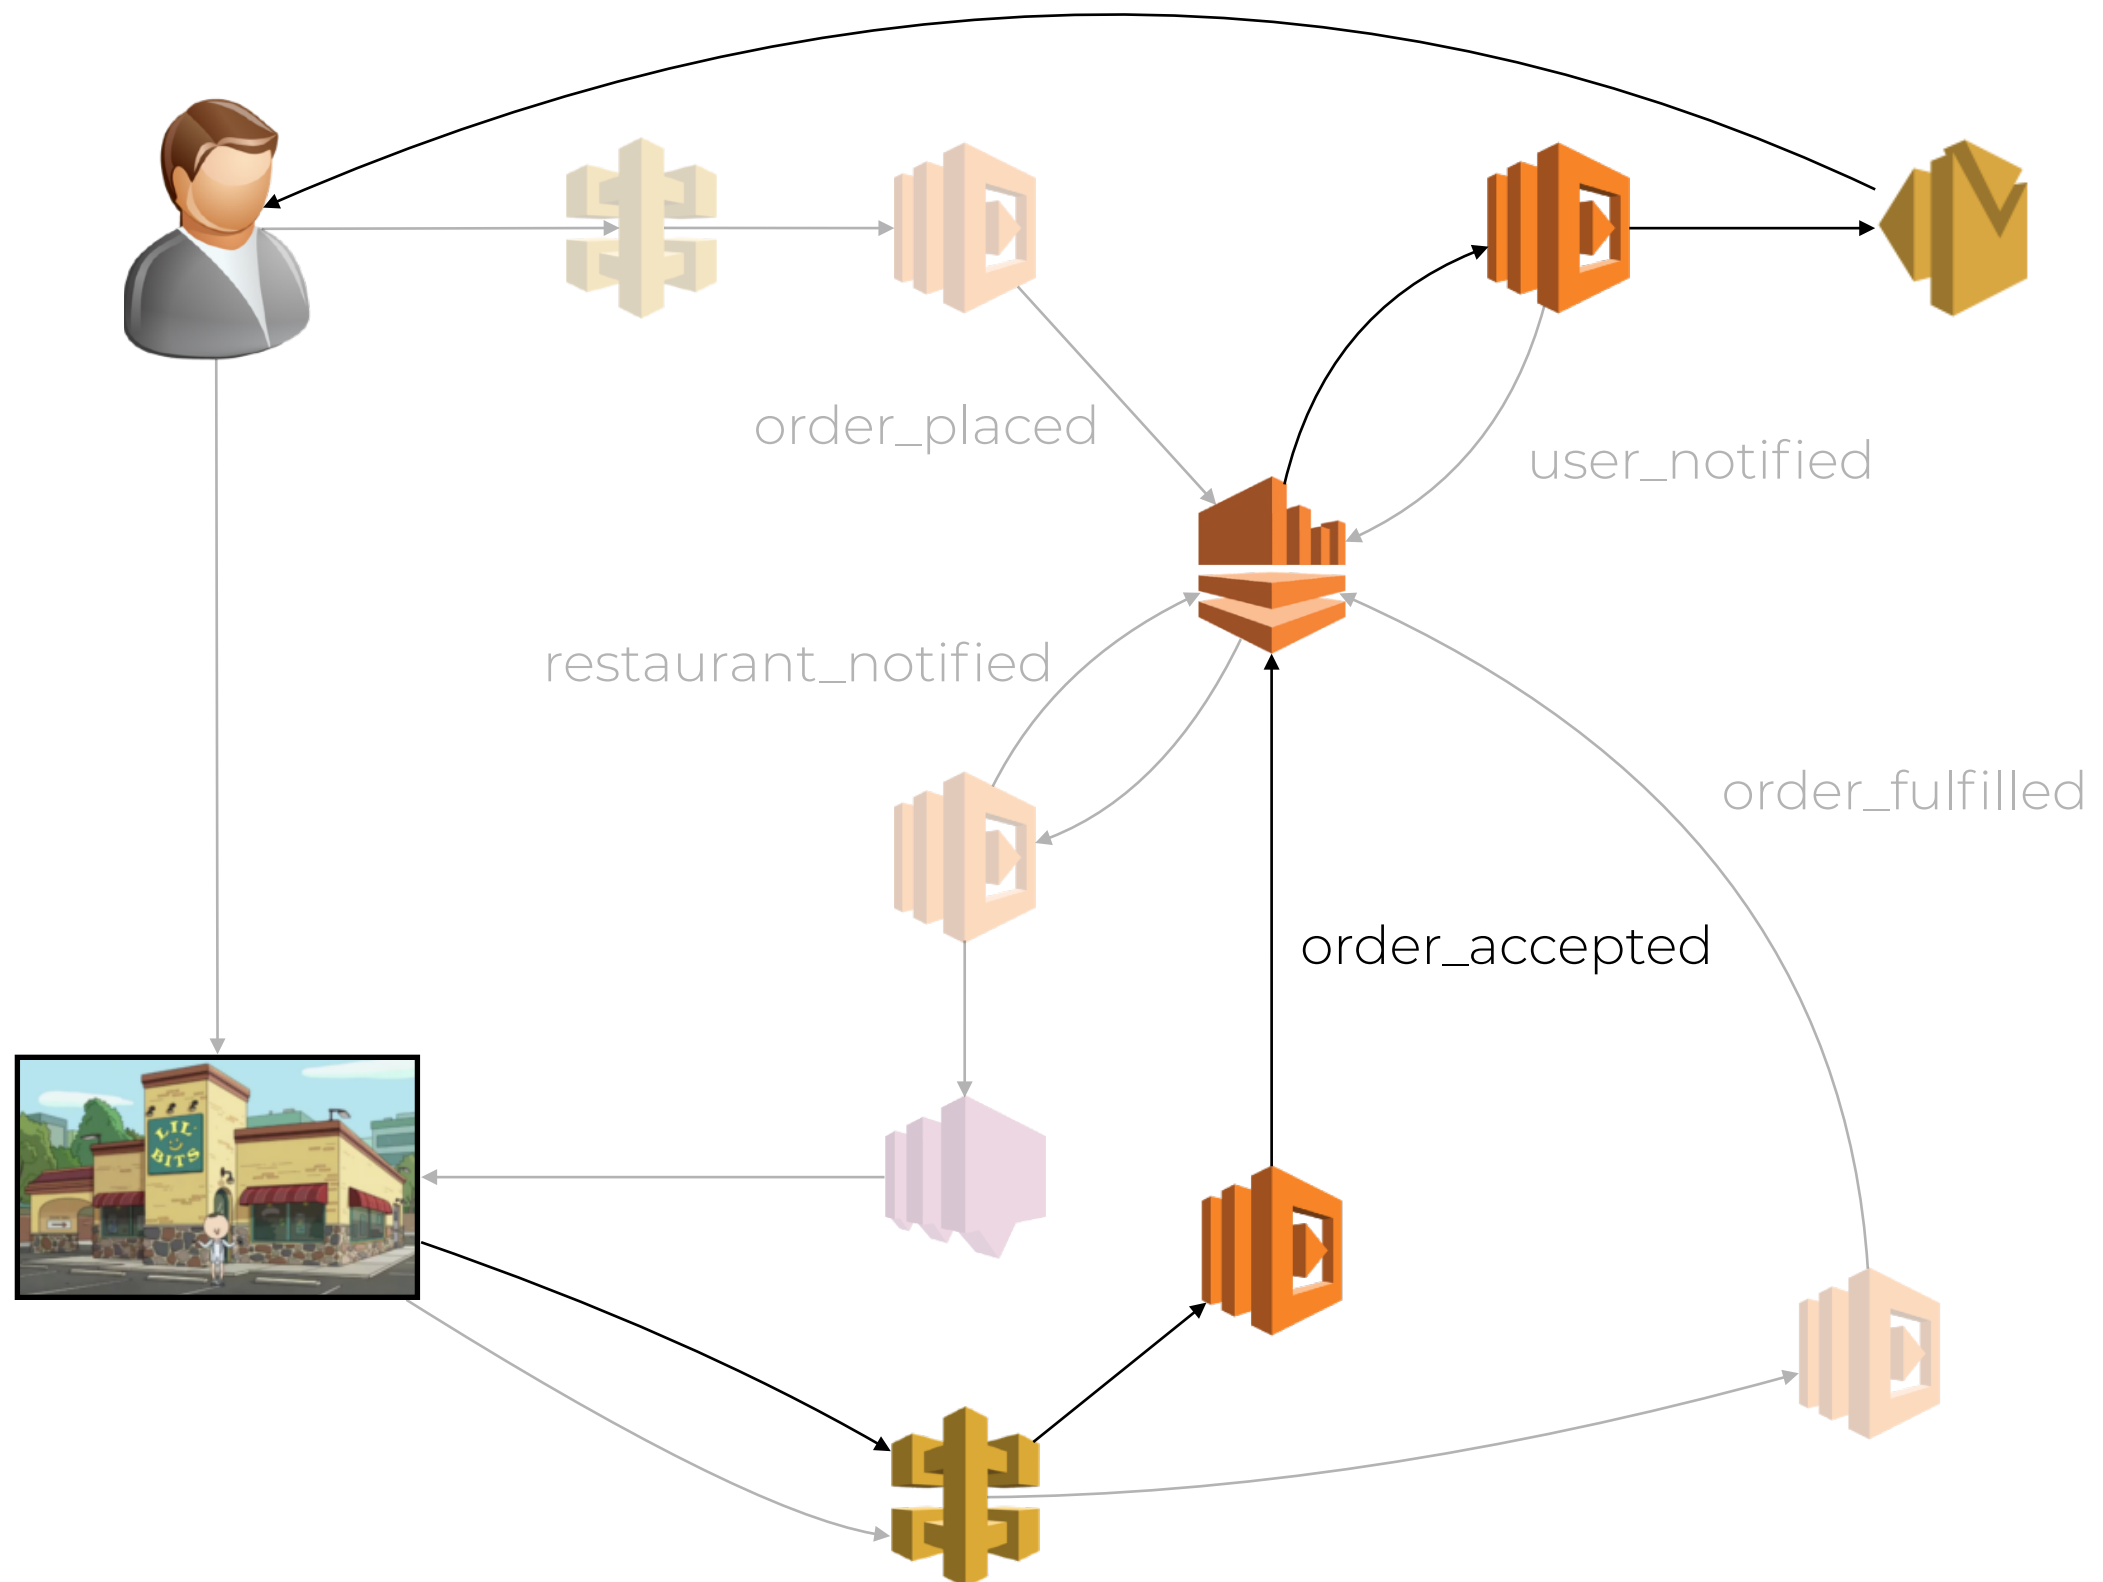


1. customer places order
2. notifies restaurant
3. restaurant accepts order
4. customer confirmation
5. pick up

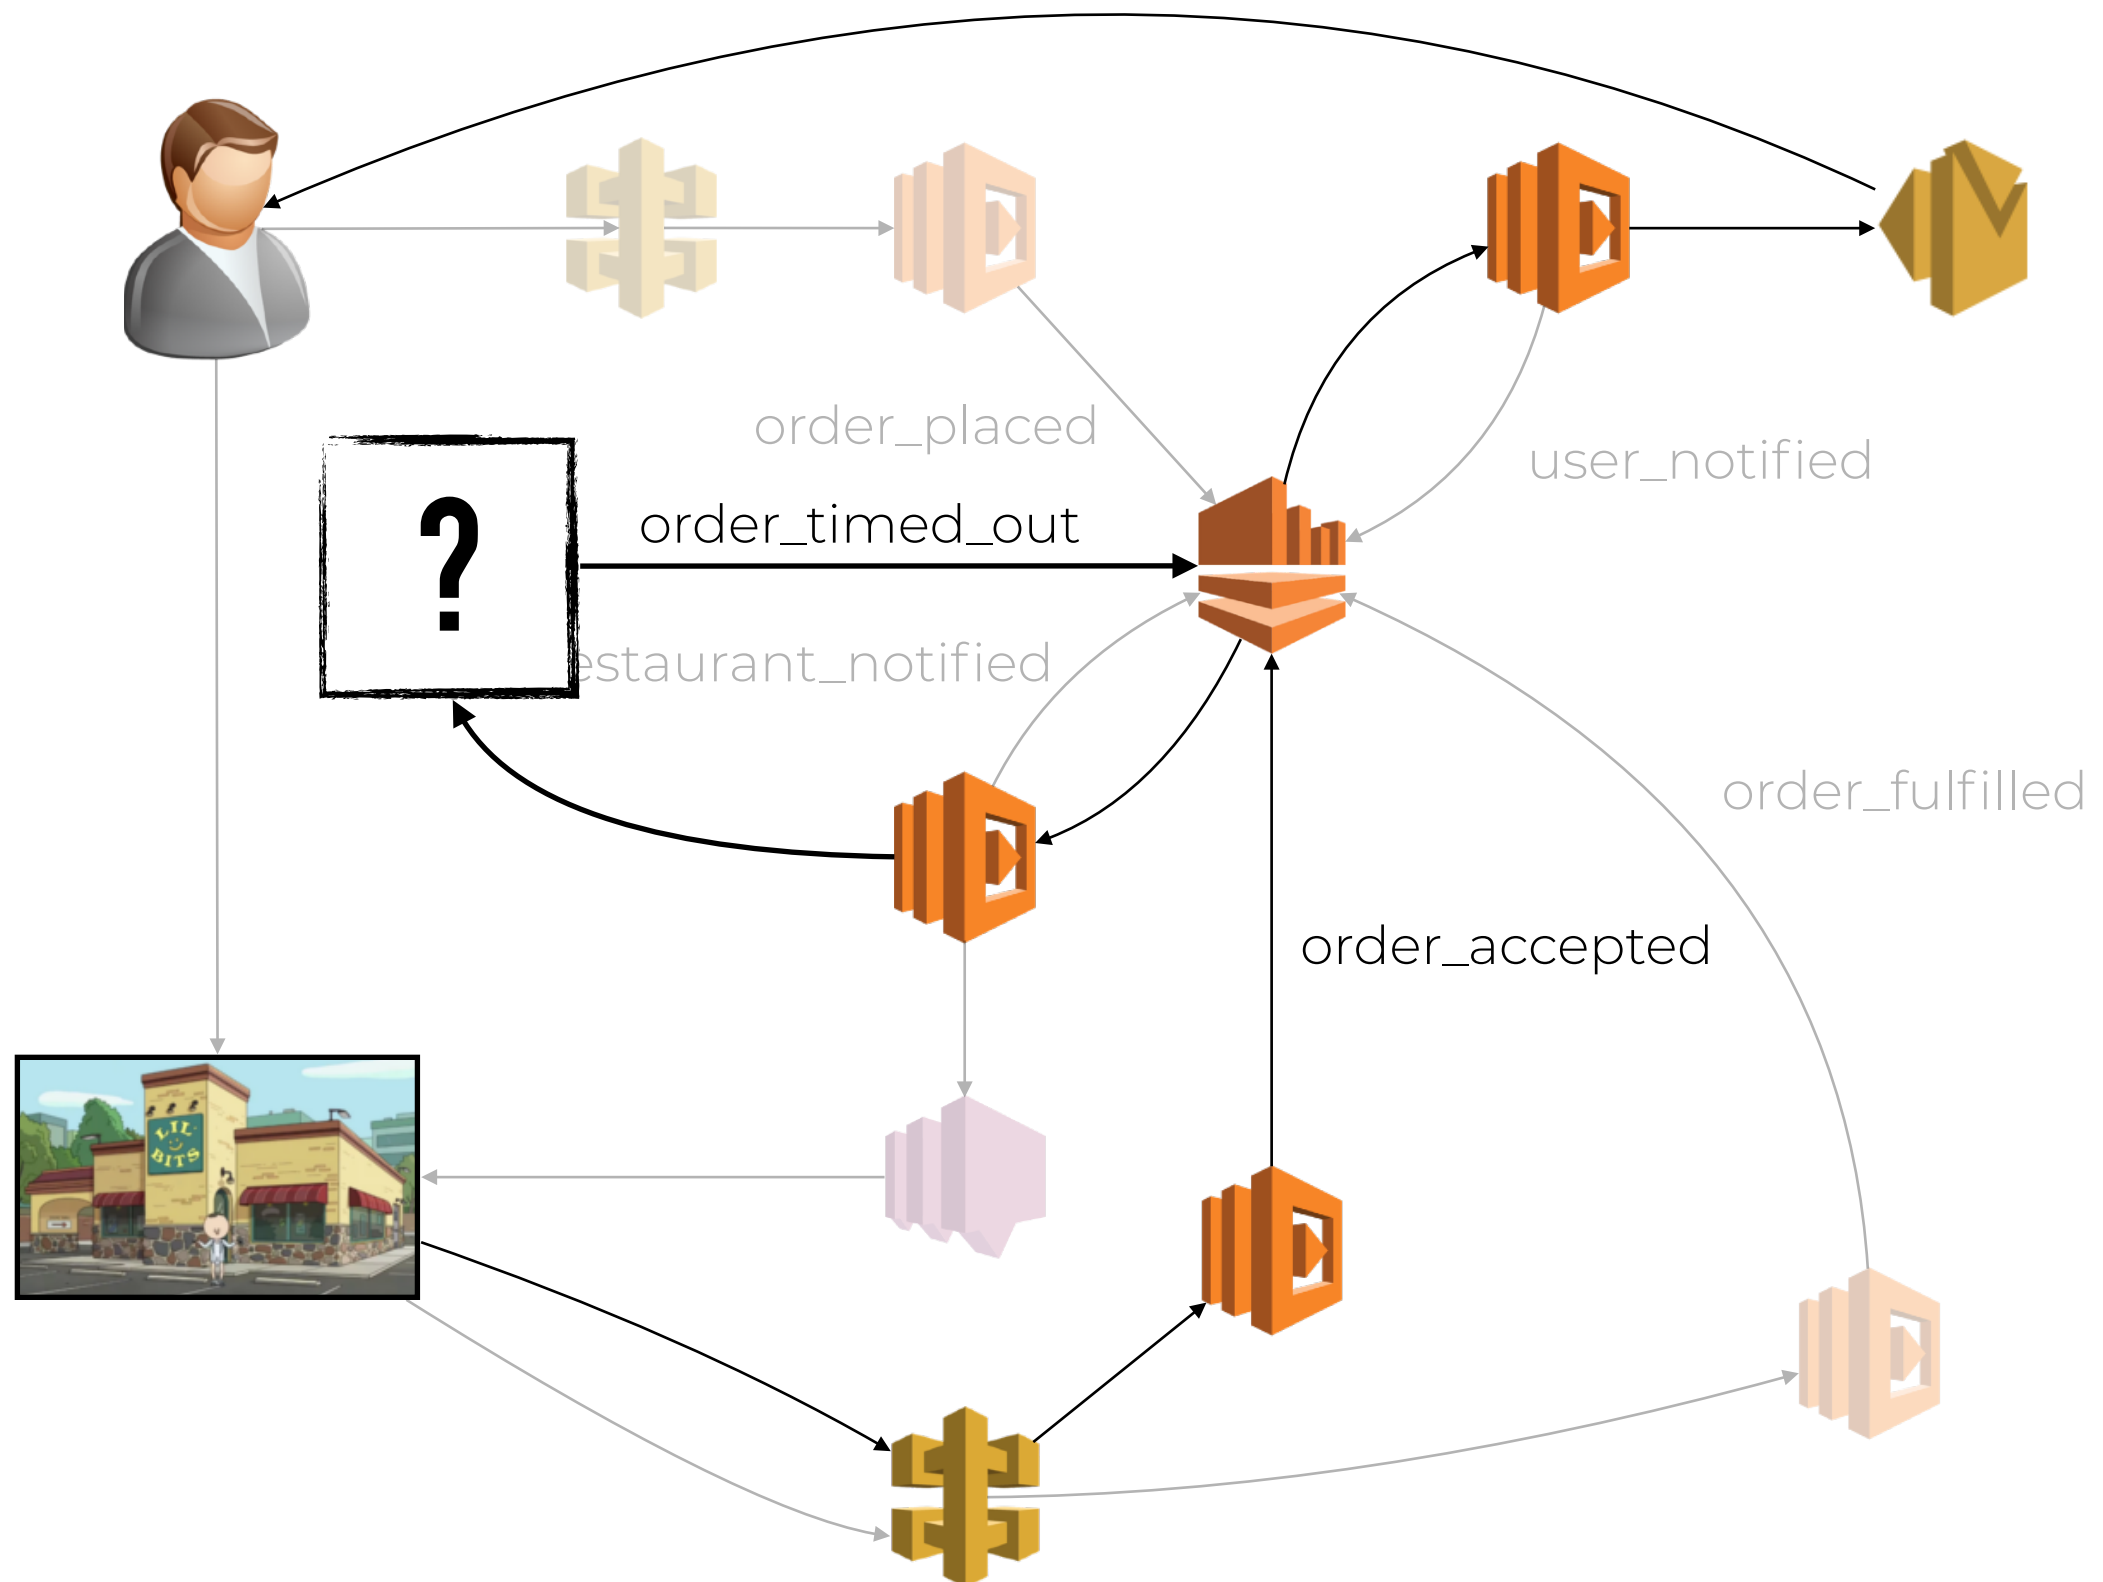
order flow



order flow



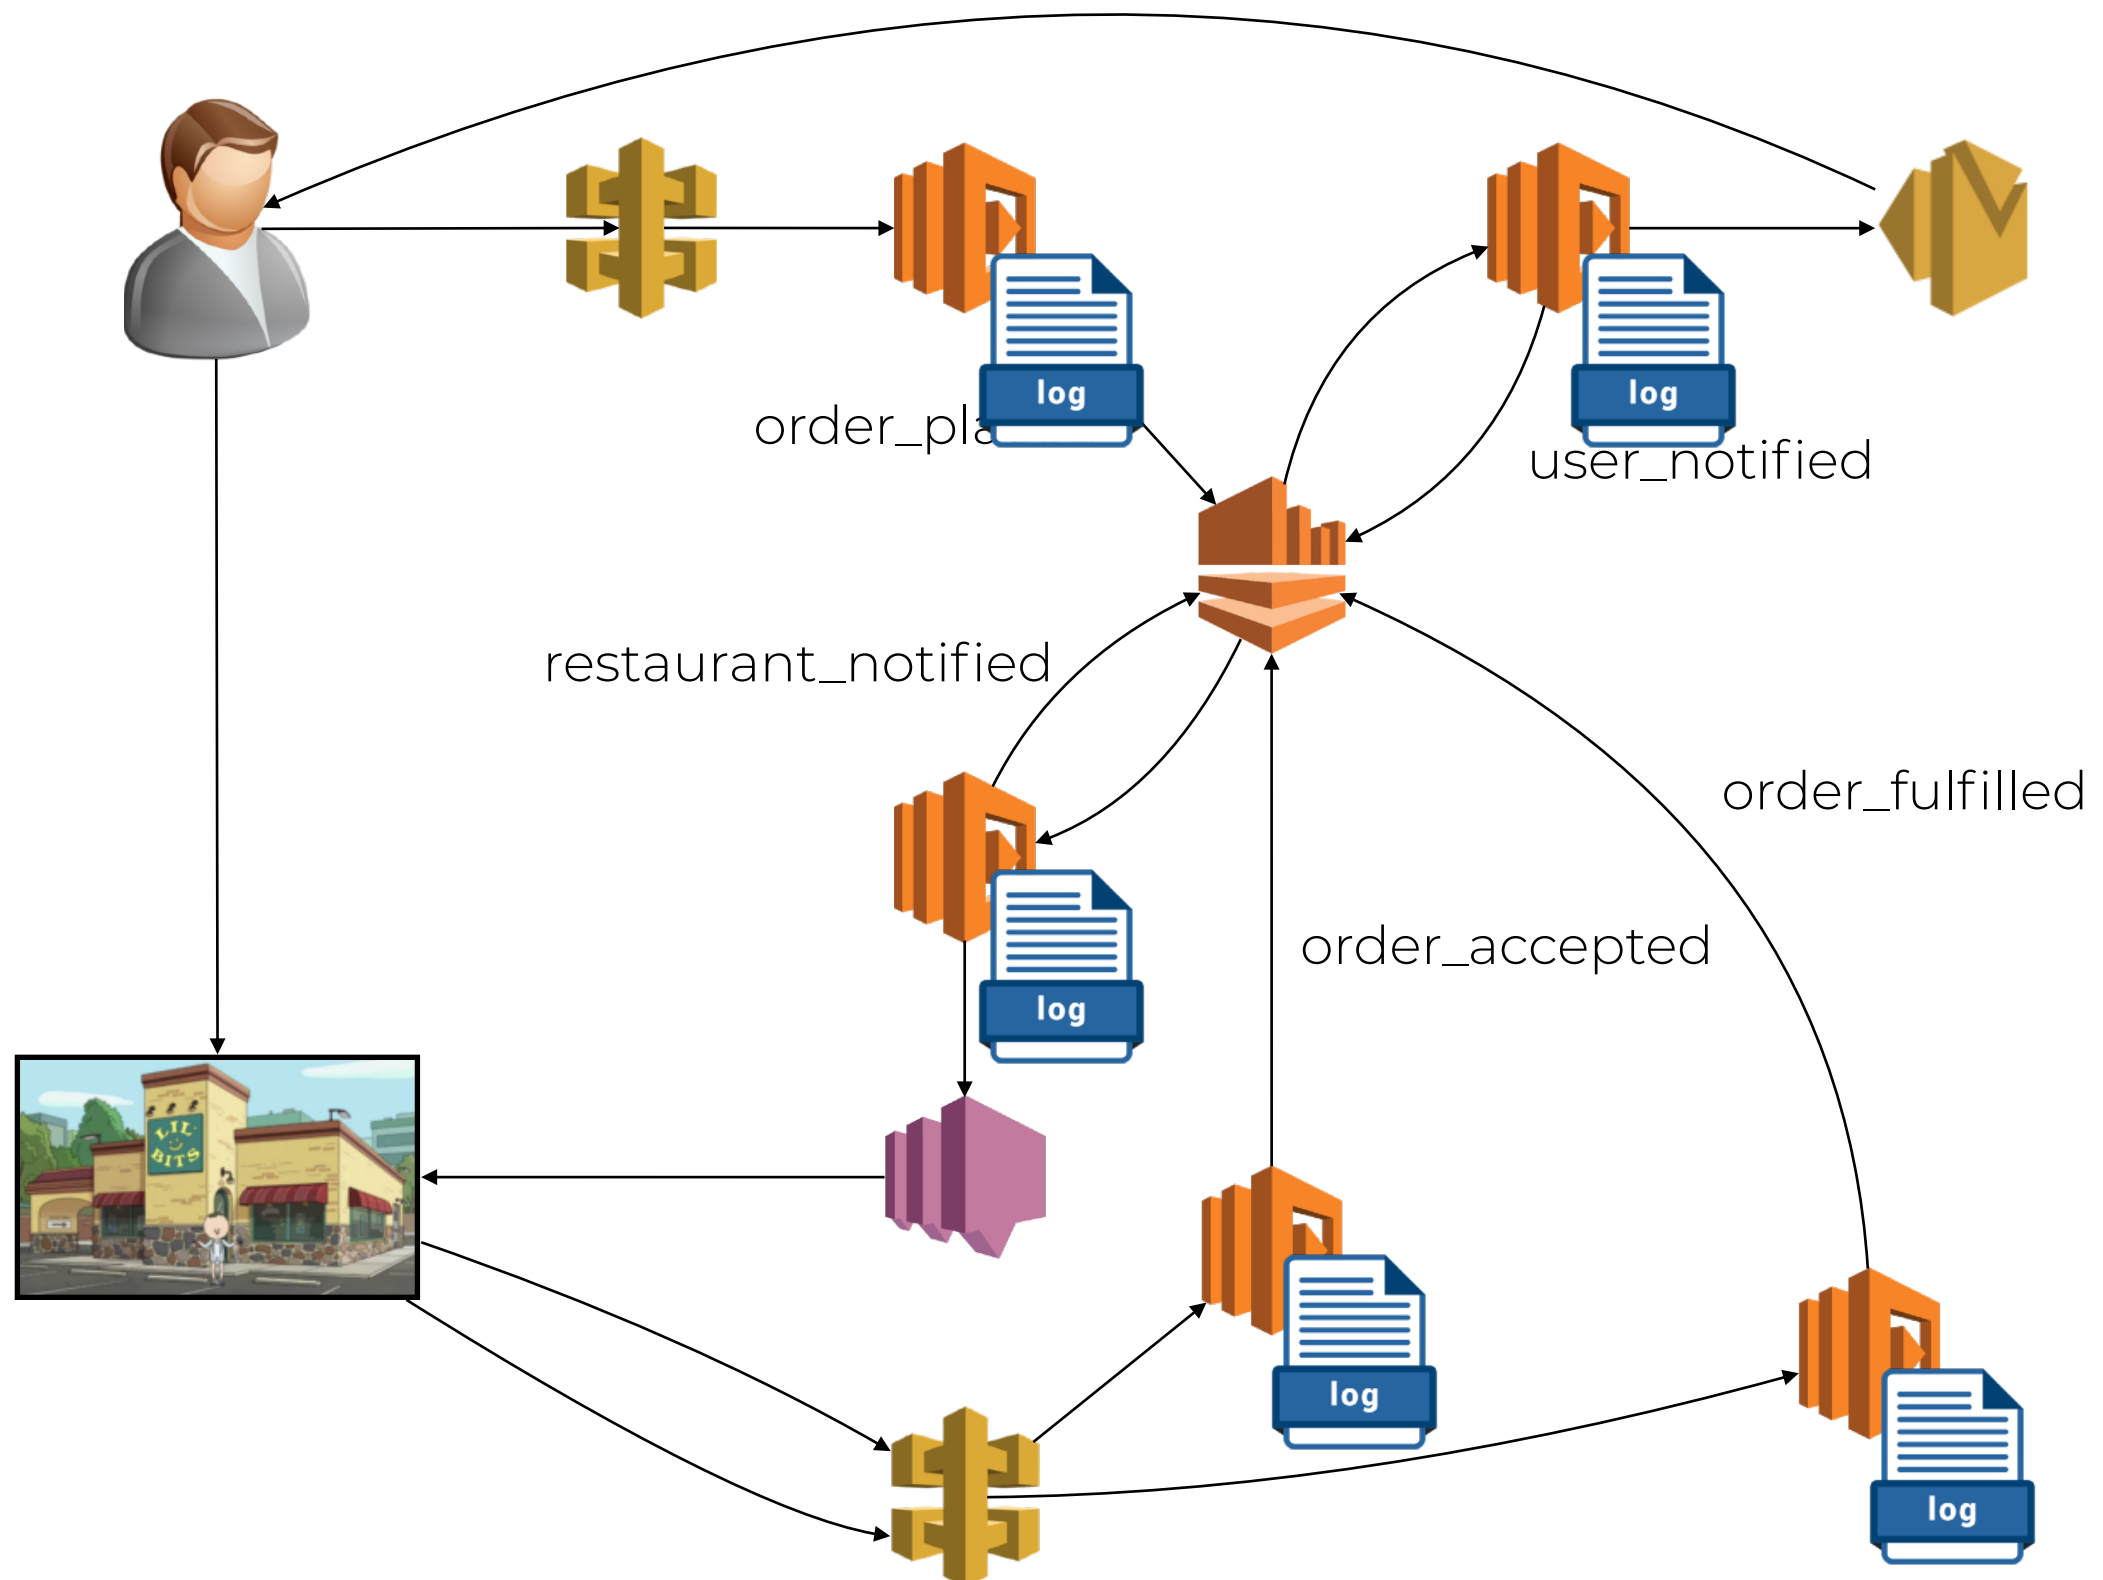
order flow



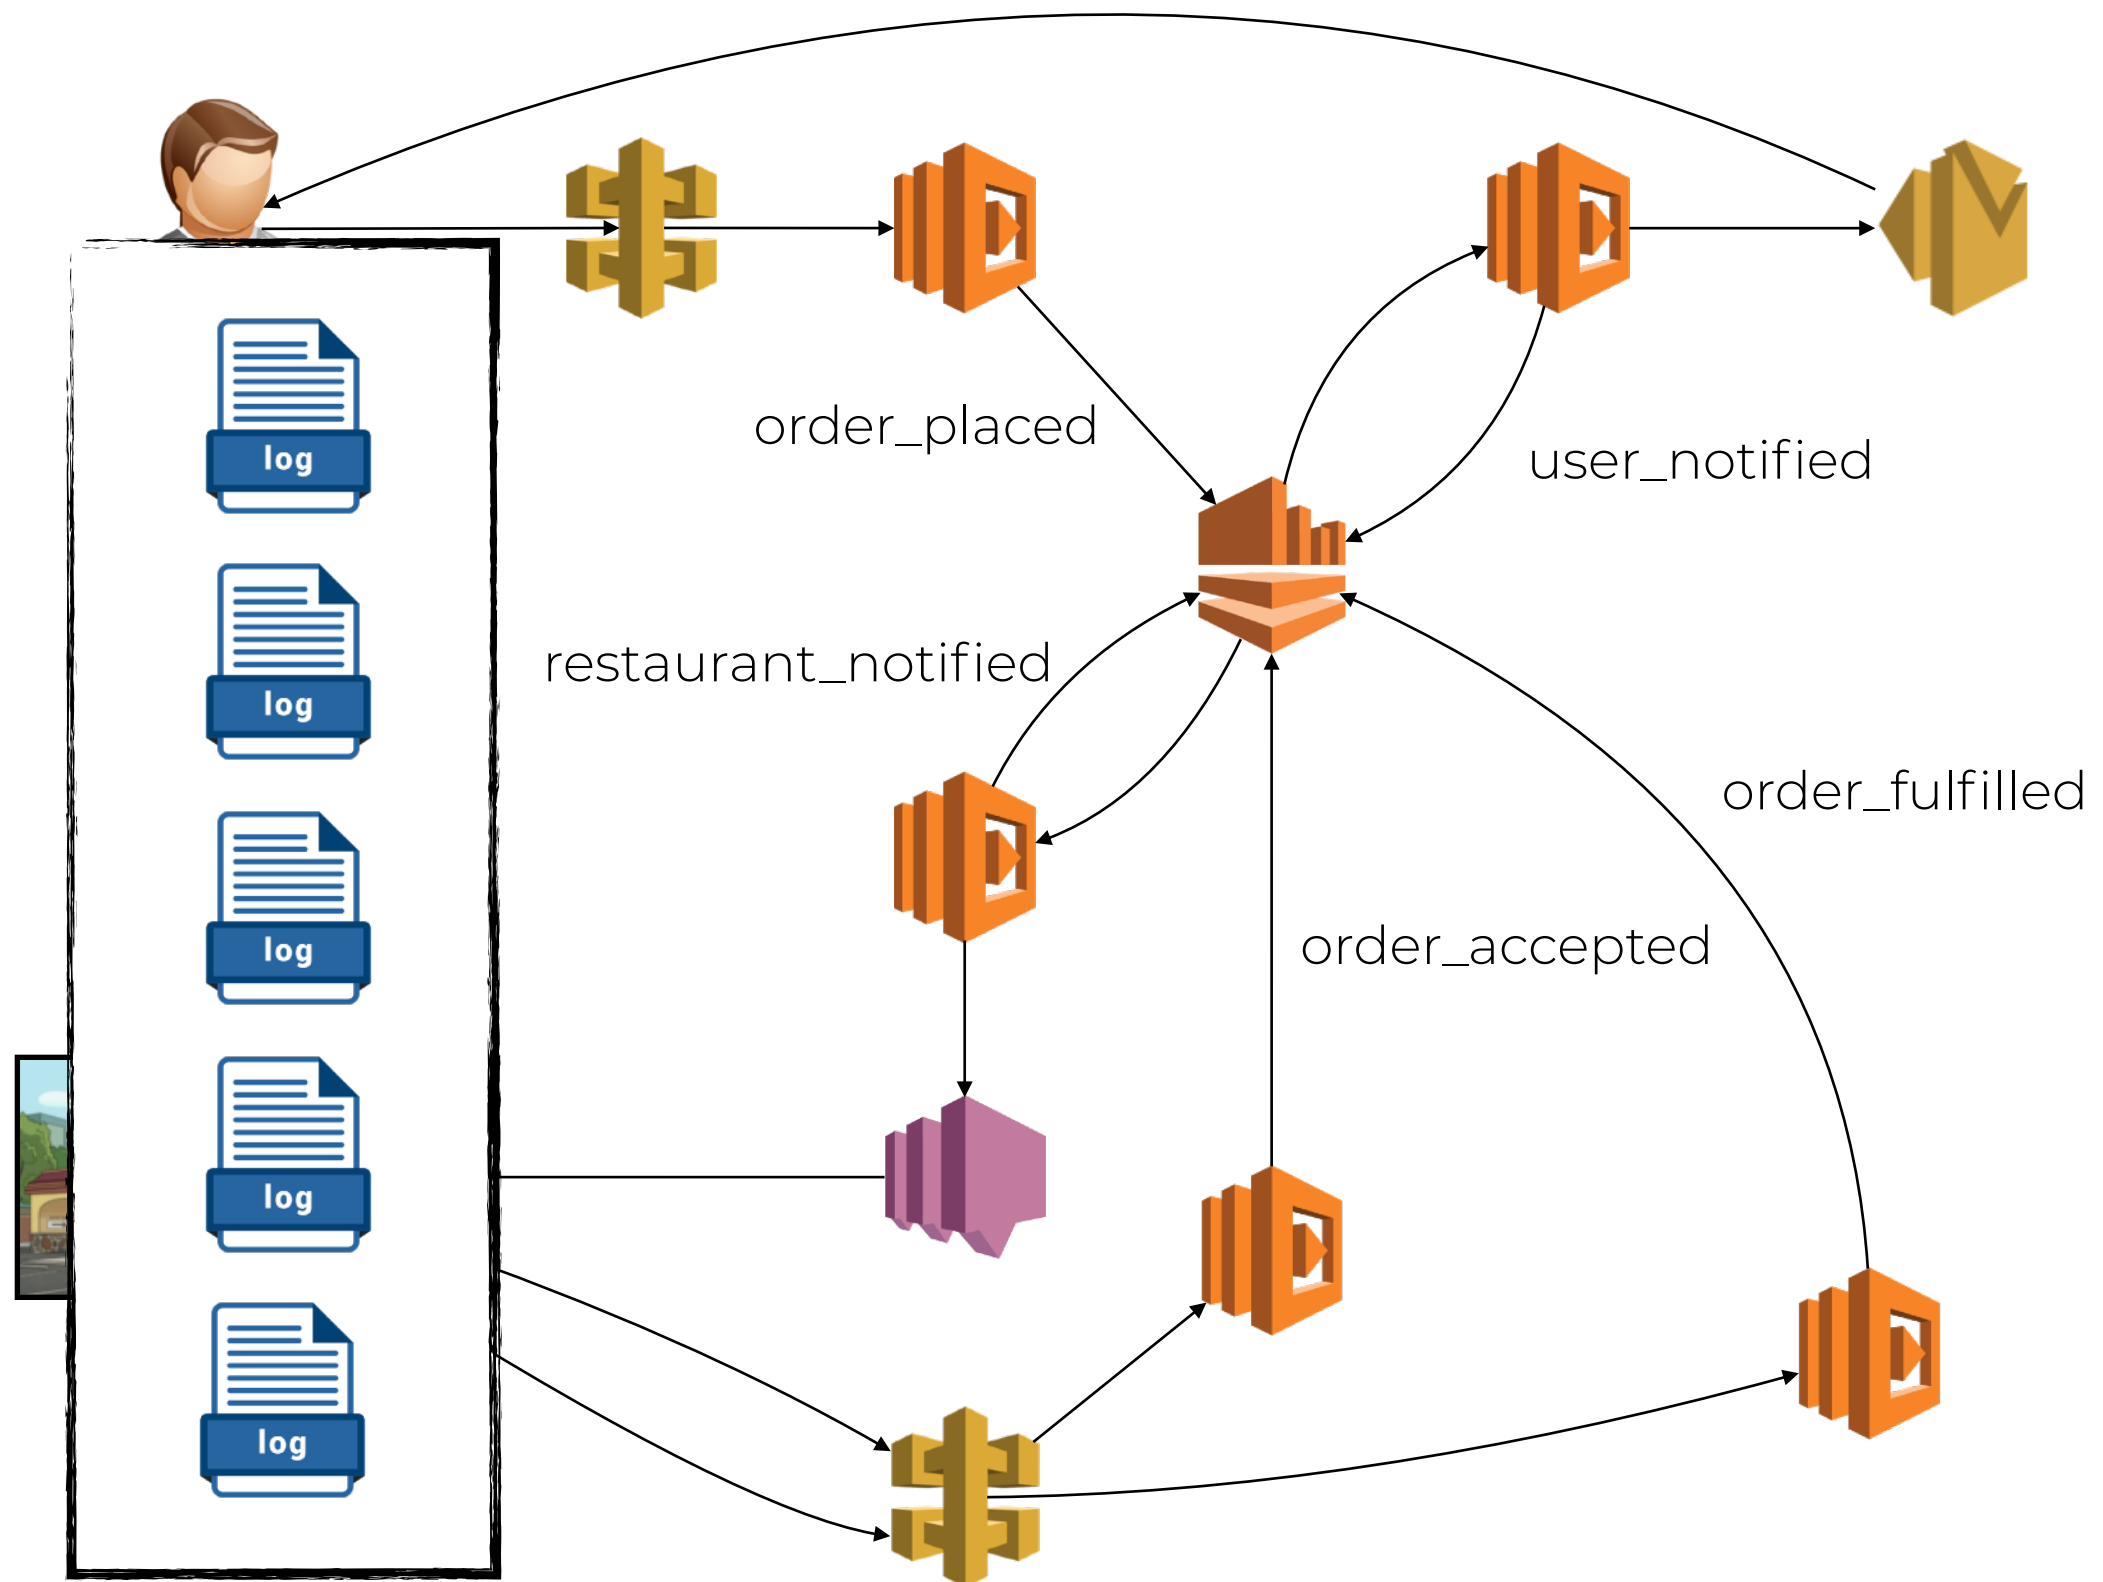
failure modes

- minimise impact on user experience
- strong traceability

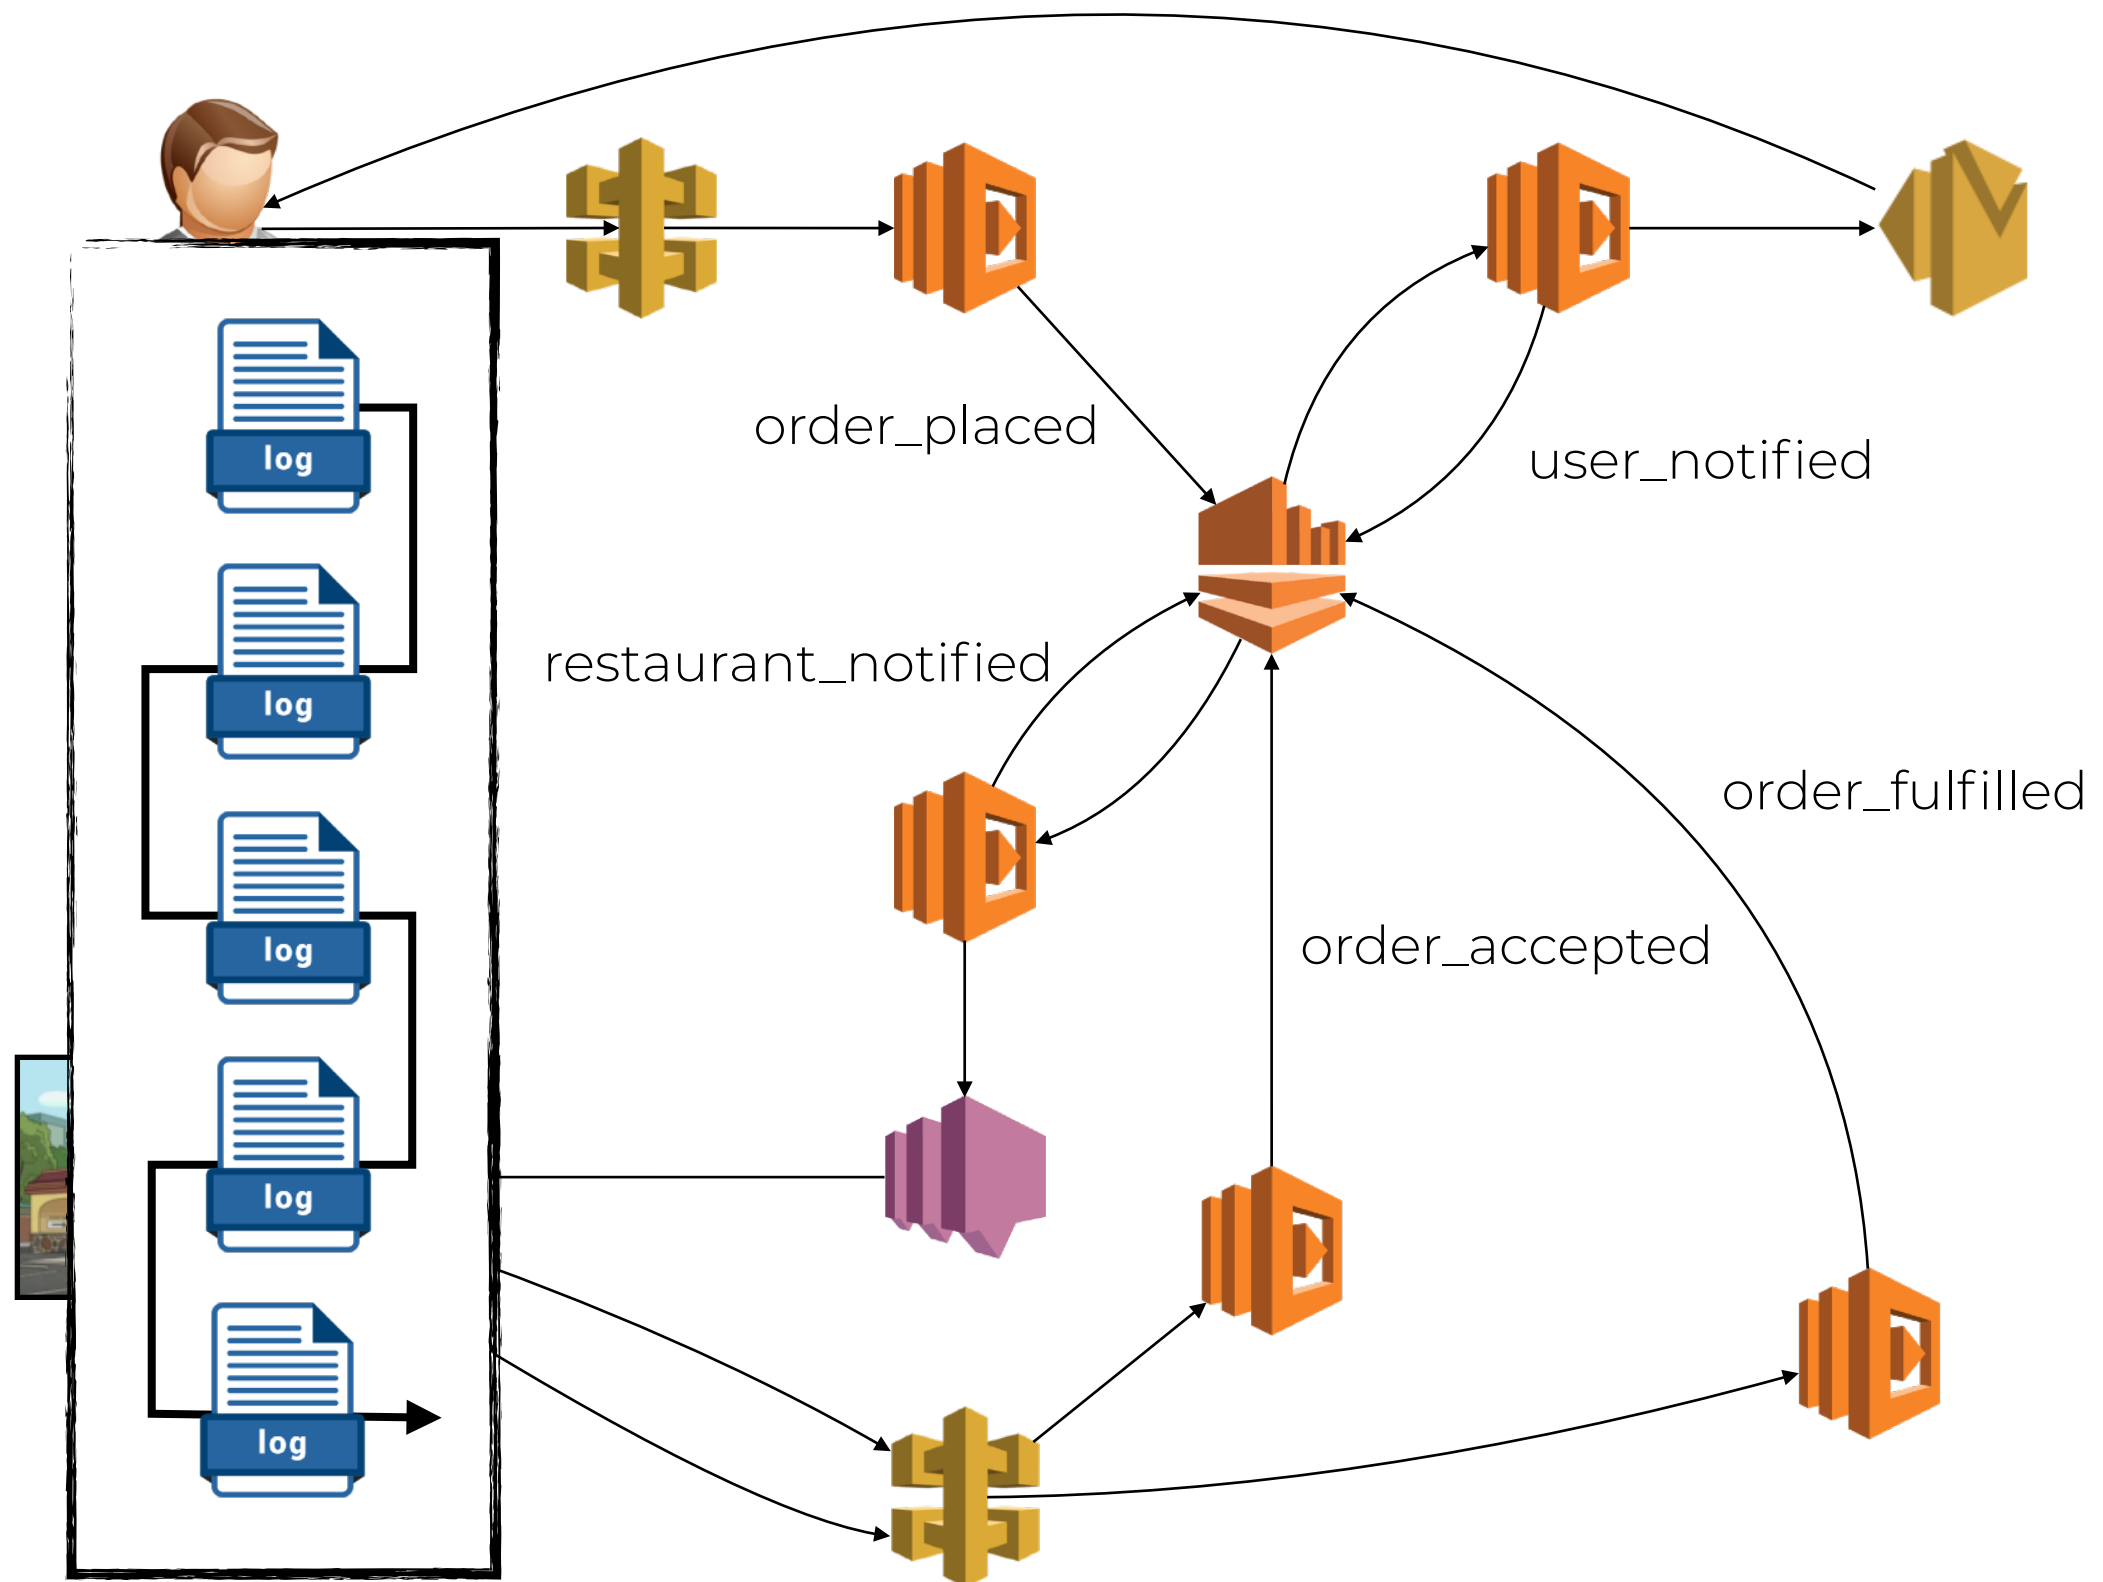
order flow



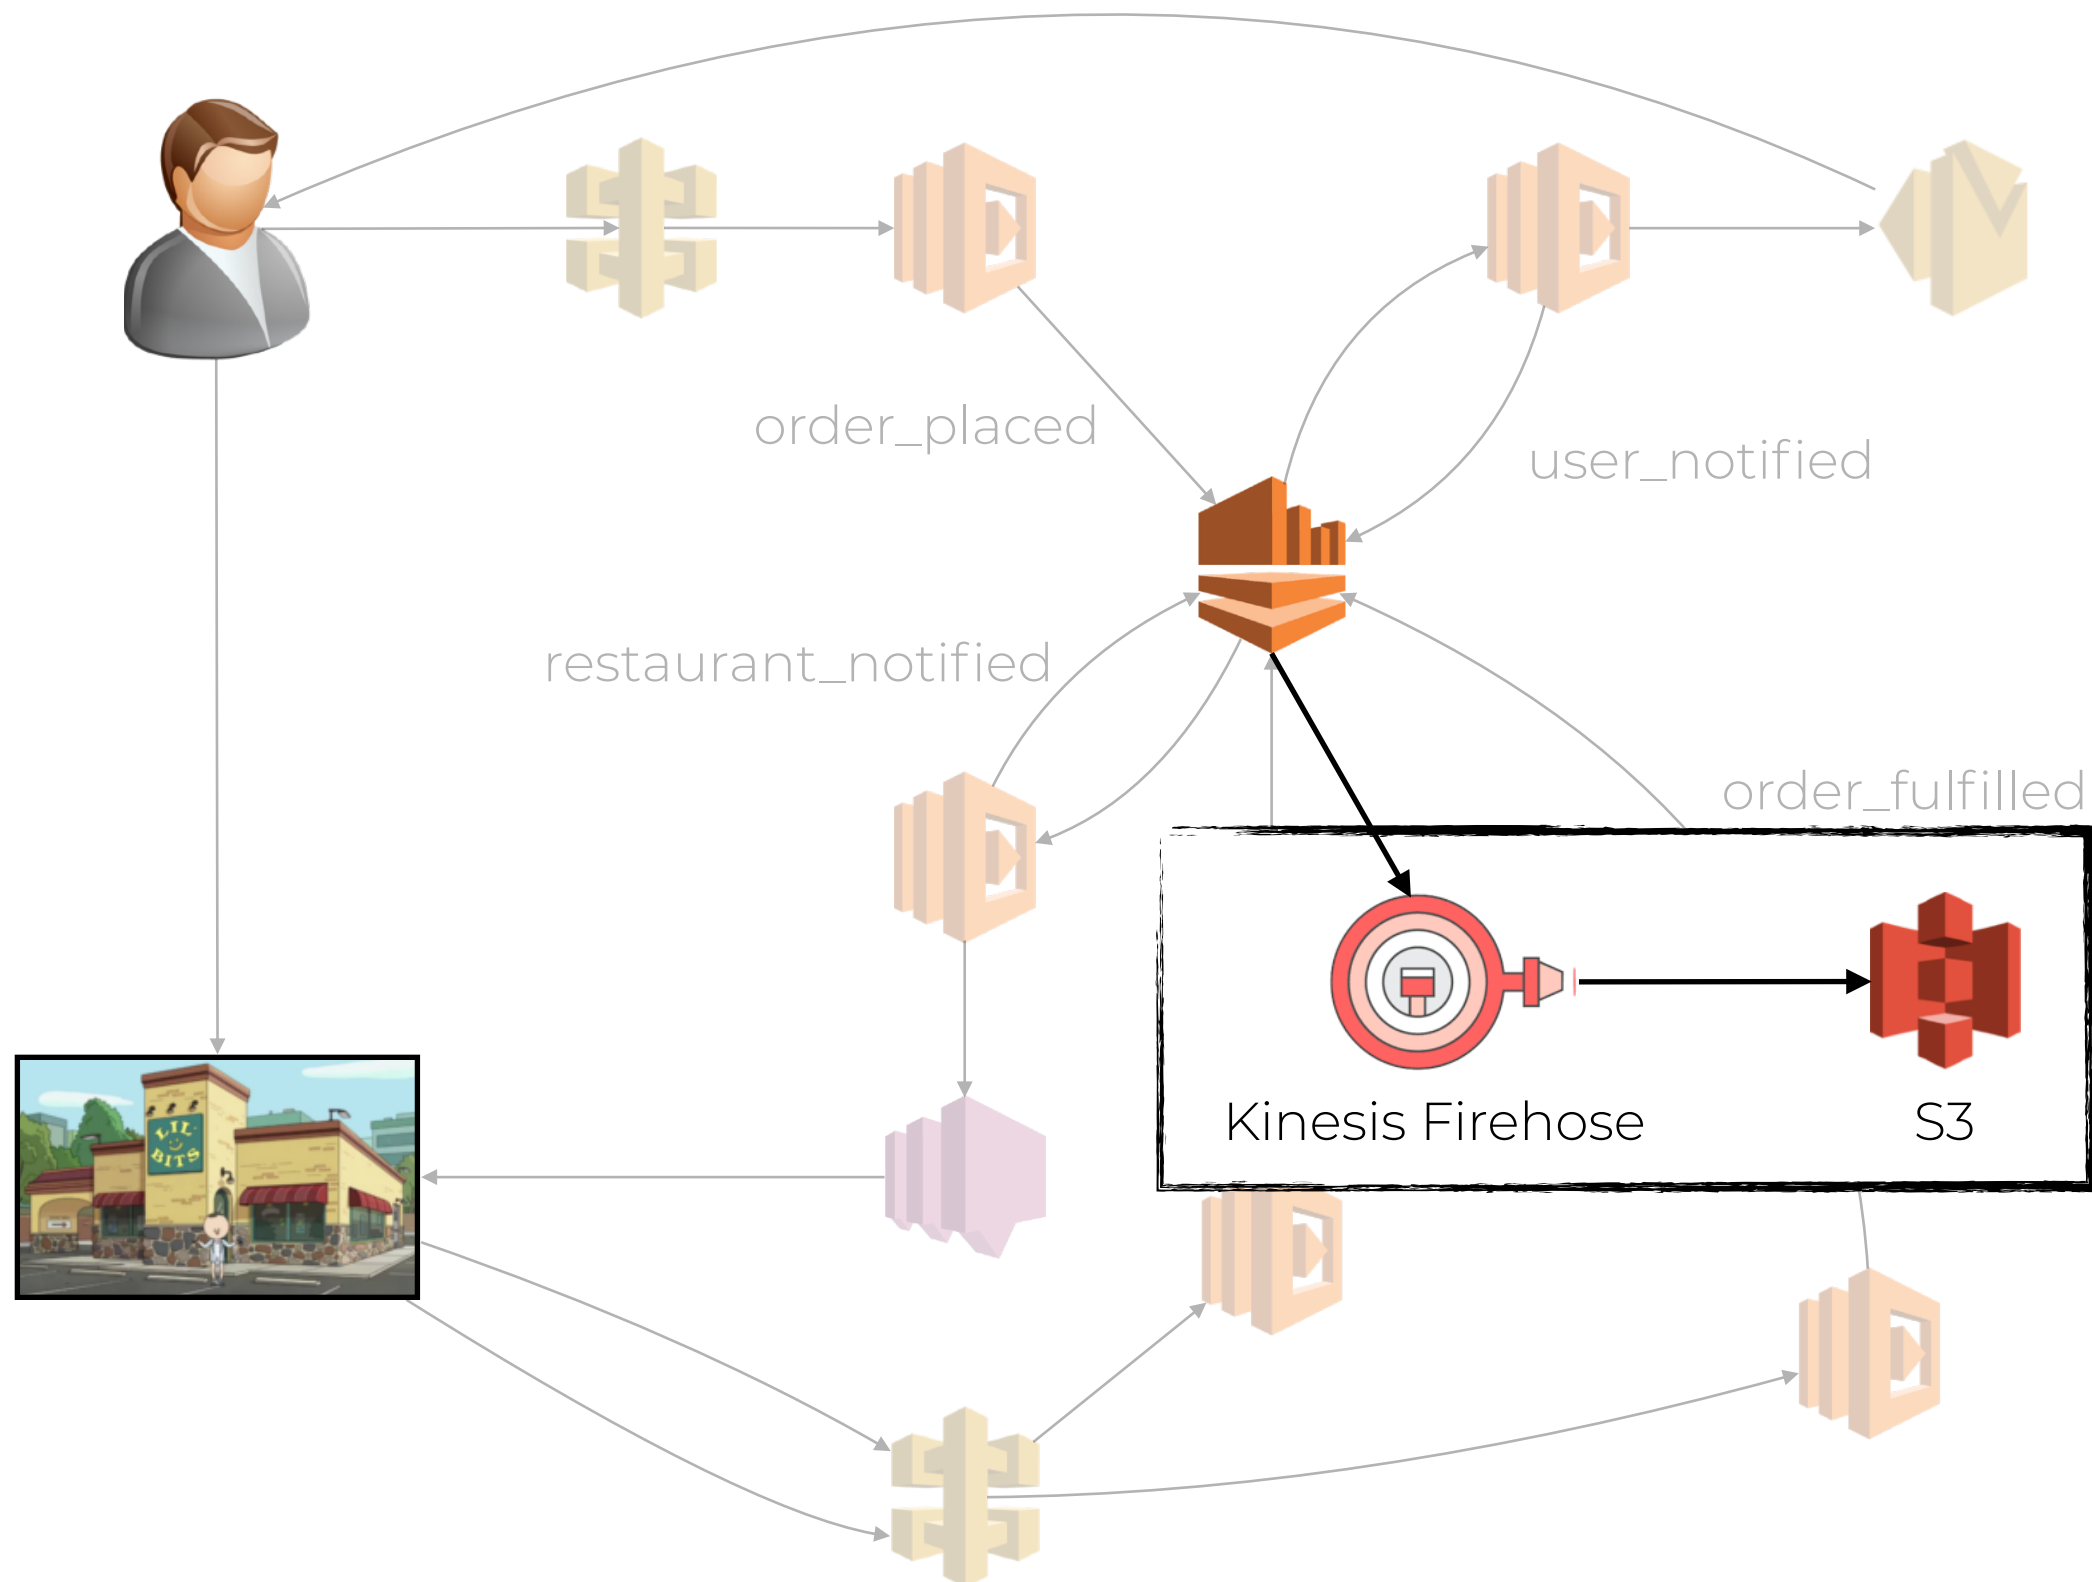
order flow



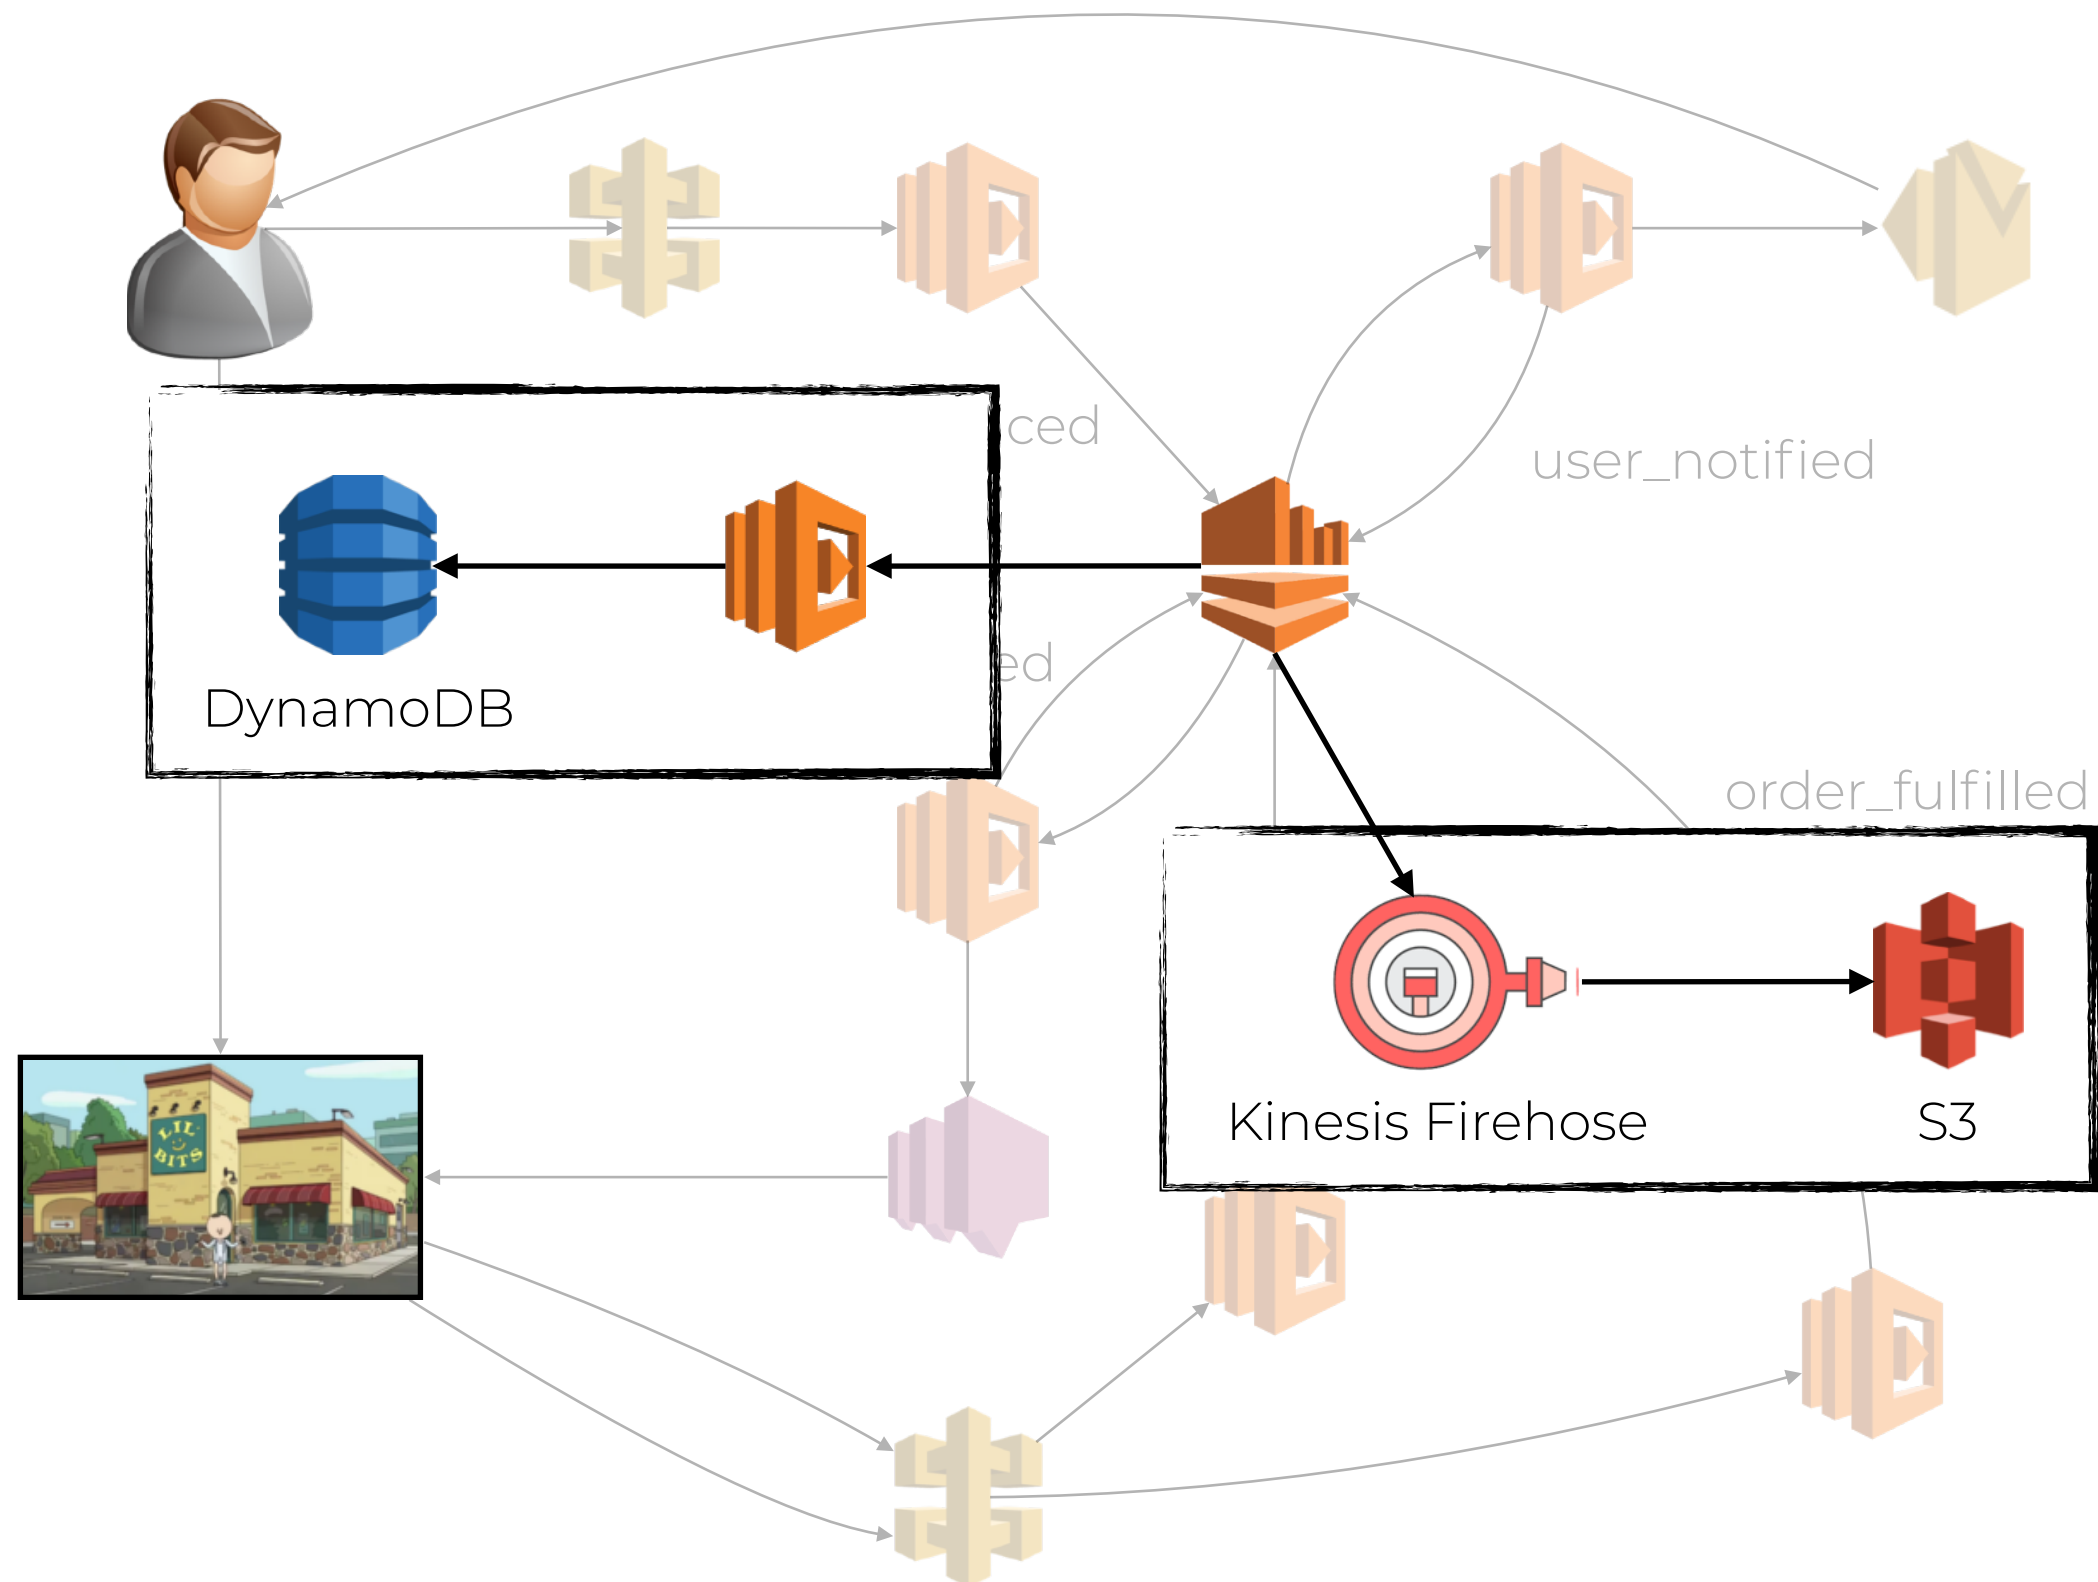
order flow



order flow




order flow






AWS Step Functions

order flow



order flow



order flow

